

Boston House Price Prediction

In this project, I have built a machine learning model to predict house prices using two different algorithms:

1. Linear Regression
2. XGBoost Regression

Since this is a **regression problem**, the task involves predicting continuous numerical values (house prices) based on input features like the number of rooms, crime rate, and more. This approach allows us to understand how these features influence house prices and make accurate predictions for unseen data.

In [43]:

```
#IMPORTING NECESSARY LIBRARIES
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
from sklearn.metrics import mean_squared_error, r2_score
from xgboost import XGBRegressor
```

In [3]:

```
#IMPORTING DATASET
House_price_dataset = pd.read_csv('/content/BostonHousing.csv')
```

In [8]:

```
print(House_price_dataset)
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	
	ptratio	b	lstat	medv							
0	15.3	396.90	4.98	24.0							
1	17.8	396.90	9.14	21.6							
2	17.8	392.83	4.03	34.7							
3	18.7	394.63	2.94	33.4							
4	18.7	396.90	5.33	36.2							
..							
501	21.0	391.99	9.67	22.4							
502	21.0	396.90	9.08	20.6							
503	21.0	396.90	5.64	23.9							
504	21.0	393.45	6.48	22.0							
505	21.0	396.90	7.88	11.9							

[506 rows x 14 columns]

In [6]:

```
#DISPLAY FIRST FEW ROWS TO VERIFY THE IMPORT
print(House_price_dataset.head())
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	ls
tat medv													
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4
.98	24.0												
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9
.14	21.6												
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4
.03	34.7												
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2
.94	33.4												
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5
.33	36.2												

This dataset contains **506 records** and **14 features**, representing various characteristics of homes located in Boston. Each feature provides specific information, such as the number of rooms, crime rate, and property tax, which can help predict house prices in the area.

Steps of Steps by Step Approach in the Project:

- 1. Collecting Data**
The first step is gathering the dataset, which contains information about various homes in Boston, such as features like the number of rooms, crime rates, etc., that will help us predict house prices.
- 2. Pre-processing Data**
The data is not yet ready to be fed into a machine learning model. Pre-processing is required to clean and format the data, handle missing values, and scale the features so the model can process them effectively.
- 3. Exploratory Data Analysis (EDA)**
In this step, we analyze the dataset to understand the relationships between features. For example, we find correlations between different features (like the number of rooms or crime rates) to see which ones are most closely related to the target variable (house prices).
- 4. Splitting Data into Train-Test Split**
The data is divided into two sets: one for training the model (training set) and the other for testing the model's performance (test set). This helps ensure the model generalizes well to unseen data.
- 5. Feeding the Data to XGBoost Regressor for Prediction**
The training data is used to train the **XGBoost Regressor** model, which makes predictions on the test data. This model is known for its efficiency and accuracy in regression tasks.
- 6. Evaluation**
Finally, the model's performance is evaluated by comparing its predictions with the actual values using metrics like Mean Squared Error (MSE) and R-squared (R²). This step helps us assess how well the model is performing and whether it is a good fit for the problem.

In [7]:

```
#DISPLAY LAST FEW ROWS
print(House_price_dataset.tail())
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	l
stat medv													
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	
9.67	22.4												
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	
9.08	20.6												
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	
5.64	23.9												
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	
6.48	22.0												
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	
7.88	11.9												

In [9]:

```
House_price_dataset.shape
```

Out[9]:

(506, 14)

In [10]:

```
House_price_dataset.isnull().sum()
```

Out[10]:

	0
crim	0
zn	0
indus	0
chas	0
nox	0
rm	0
age	0
dis	0
rad	0
tax	0
ptratio	0
b	0
lstat	0
medv	0

dtype: int64

In [12]:

```
House_price_dataset.describe()
```

Out[12]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	t
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.2371
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.5371
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.0000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.0000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.0000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.0000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.0000

In [18]:

```
correlation = House_price_dataset.corr()
```

In [21]:

```
# constructing a heatmap to understand the correlation
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.2f', annot=True, annot_kws={'size':8}, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show
```


Out[21]:

```
matplotlib.pyplot.show
def show(*args, **kwargs)
```

Display all open figures.

Parameters

block : bool, optional



Splitting the data into train_test_split

In [32]:

```
# Separate features and target variable
X = House_price_dataset.drop(columns=['medv']) # Replace 'medv' with the correct target
column name in your dataset
Y = House_price_dataset['medv']
```

In [29]:

```
print(X)
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	b	lstat
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
..
501	21.0	391.99	9.67
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

Note:

The target variable, which represents house prices, was previously named 'MEDV' in the dataset. Additionally, the prices are in thousands, meaning that the values in the dataset represent the house prices in thousands of dollars. For example, a predicted price of 25 represents 25,000 dollars. Always remember to multiply by 1,000 when interpreting the actual house price.

In [33]:

```
print(Y)
```

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501     22.4
502     20.6
503     23.9
504     22.0
505     11.9
Name: medv, Length: 506, dtype: float64
```

Splitting the data into train test split

In [34]:

```
X_train , X_test , Y_train , Y_test =train_test_split(X,Y,test_size=0.2,random_state=42)
```

In [36]:

```
print(X.shape,X_train.shape,X_test.shape)
```

```
(506, 13) (404, 13) (102, 13)
```

Standardize the data (optional but recommended for models sensitive to scale)

Standardizing the data is an important step, especially for models like Linear Regression and XGBoost that are sensitive to the scale of the features. Standardization ensures that each feature has a mean of 0 and a standard deviation of 1, which helps the model converge faster and improves performance.

Why standardizing?

Some machine learning models (like Linear Regression) are sensitive to the scale of the data. If the features vary widely (e.g., one feature ranges from 1 to 10 and another ranges from 1,000 to 100,000), the model might struggle to learn effectively. By standardizing the data, all features will contribute equally, allowing the model to learn more efficiently.

In [39]:

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

now we Train the machine learning model by using 1)linear regression and then by 2)by xgboost regressor

1)By using Linear Regression

In [46]:

```
model = LinearRegression()
model.fit(X_train_scaled, Y_train)
```

Out[46]:

▼ LinearRegression i ?

LinearRegression()

In [47]:

```
# Making predictions
y_pred = model.predict(X_test_scaled)
```

In [48]:

```
# Evaluating the model
from sklearn.metrics import mean_squared_error, r2_score
```

In [53]:

```
mse = mean_squared_error(Y_test, y_pred)
r2 = r2_score(Y_test, y_pred)


print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R-squared (R2 Score): {r2:.2f}")
```


Mean Squared Error (MSE): 24.29
R-squared (R2 Score): 0.67

In [55]:

```
#Visualize Predictions
plt.figure(figsize=(8, 6))
plt.scatter(Y_test, y_pred, alpha=0.6, color='blue')
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], color='red', linestyle='--')
plt.title("Actual vs Predicted Prices")
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.show()
```

Actual vs Predicted Prices





now we train model by using xgboost regressor

In [56]:

```
# loading the model
model = XGBRegressor()
```

In [57]:

```
# training the model with X_train
model.fit(X_train, Y_train)
```

Out[57]:

```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,

```

evaluating the model by making prediction on training data

In [58]:

```
# accuracy for prediction on training data
training_data_prediction = model.predict(X_train)
```

In [59]:

```
print(training_data_prediction)
```

```

[11.990929  19.915493  19.392988  13.408072  18.19098  24.603947
 21.08647   24.697266   8.70618   27.501347  20.708258  36.17404
 31.59529  11.69739   39.802494  13.893334  21.796898  23.695662
 17.500210  24.400176   8.700540  10.101701  25.070050  20.410500

```

17.590319	24.409176	8.799549	19.181791	25.278059	20.410582
23.10783	37.90089	15.601625	45.400623	15.706774	22.599428
14.514692	18.697655	17.797323	16.117708	20.609972	31.598558
29.095152	15.600668	17.563873	22.51095	19.401443	19.287243
8.4968405	20.607521	17.006351	17.093975	14.495169	49.98619
14.284735	12.609954	28.688086	21.203852	19.306932	23.089056
19.10372	25.004898	33.408142	4.992768	29.599434	18.685545
21.707012	23.096918	22.802471	20.991085	48.796387	14.627051
16.613852	27.075224	20.087227	19.794464	20.992268	41.290615
23.175125	20.378569	18.558722	29.399414	36.40198	24.388987
11.816533	13.792832	12.272116	17.794773	33.087368	26.73878
13.393904	14.386449	50.004574	21.98985	19.906004	23.784376
17.516703	12.68493	5.6061826	31.099165	26.202074	19.381464
16.696205	13.794351	22.894213	15.30921	27.500492	36.09182
22.87913	24.495989	24.998743	49.99985	34.889656	31.717129
24.095907	22.0973	14.106711	42.799473	19.324936	32.193943
26.385065	21.79409	21.668243	8.311317	46.69916	43.12514
31.4751	10.497389	16.68972	20.004532	33.30795	17.790648
50.000507	20.497787	23.19068	13.075348	19.610313	22.796833
28.680086	30.684517	22.906889	21.899395	23.89475	32.695194
24.292328	21.48476	24.578018	8.519734	26.393446	23.08319
15.00561	8.802389	19.371504	23.89509	24.67069	19.78167
23.754805	13.284227	28.994848	27.118698	34.574432	13.301209
15.598929	12.511904	14.590067	10.98594	24.80251	17.29882
8.103687	21.408438	15.598756	23.336	32.001125	38.690685
30.095453	20.506903	32.494488	42.29702	24.277178	20.604298
22.008326	18.205097	15.002868	6.3037786	20.095121	21.393785
28.409906	30.051968	20.797775	23.017813	14.3566675	11.7001
37.290714	17.108719	10.401778	22.964136	22.683504	20.312103
21.672386	49.99456	8.395799	18.804155	37.214493	16.09777
16.501677	22.201307	20.606125	13.506562	48.31008	23.785505
22.68971	17.39483	30.310457	36.009197	41.7029	18.303009
21.996626	18.597338	44.81159	11.908145	18.710653	16.18354
22.008686	7.1969705	20.392668	13.800741	13.002582	18.343378
23.102316	21.19985	23.106937	23.499647	49.99908	26.585985
22.182547	50.000687	8.290042	23.308157	21.713062	18.965406
18.39845	17.426226	13.403094	12.083532	26.599447	21.708506
28.396055	20.5091	21.985607	13.901346	11.330144	29.907621
26.616785	10.508602	23.173222	24.391636	46.002525	21.908016
7.505983	36.194332	43.98905	17.795595	27.452951	37.594242
14.098721	28.08458	10.222061	19.139948	43.813446	27.906195
25.018572	15.995898	16.602438	13.2132845	49.993385	22.189535
32.906765	15.214496	14.794672	13.834945	24.293371	33.80743
22.299568	49.998867	9.522176	13.319194	22.216854	18.106277
17.994871	25.017178	16.499542	23.017021	20.09075	33.001865
24.81743	18.206358	13.117948	34.909748	10.209828	19.903967
27.898724	23.279646	35.09593	12.776962	22.004248	18.491093
25.118921	22.485199	22.41663	28.594748	19.526266	24.799929
24.45743	21.411762	33.110527	22.894064	20.68664	24.094328
50.00028	24.701363	28.668966	7.199975	36.96739	20.306965
30.098255	19.504751	23.379314	11.4862795	21.601028	14.903119
15.188632	19.379465	8.403532	28.002205	22.611923	13.498683
14.488013	30.9902	10.897202	21.884697	22.024958	18.993114
21.374176	25.002413	17.514051	36.499477	20.103745	20.379963
16.224405	23.605675	7.403649	35.20376	50.00256	19.299261
21.221973	15.591636	33.40674	19.125359	20.995052	23.709246
18.89621	16.813063	19.715427	17.744104	22.595695	11.790642
34.931126	20.562077	20.206211	31.961233	22.303417	23.28739
14.408953	31.195423	23.981024	29.59999	19.535425	21.596525
19.961027	26.995071	33.205605	15.399309	30.504189	7.2087
23.920792	16.292902	23.906166	49.996964	22.826248	15.398463
19.233753	19.59635	22.620607	33.200974	49.997913	22.256226
14.905368	19.808466	23.701315	18.996	20.30573	11.922065
13.598755	29.814579	21.705639	19.491856	21.109499	24.519201
13.40267	18.598032]			

Now we find r squared error and mean absolute error which are evaluation metrics. In the step below we import performance Evaluation metrics like mean_squared_error, r2_score for Model Evaluation. The two key evaluation metrics used here are mean absolute error and R-squared error. They are important for determining how well a regression model predicts target variables. In the above Model we predict yield (target variable) based on DOY (day of year). The Error Metrics are discussed below

1) Mean_squared_error : Mean Squared Error (MSE) is a measure used metric to calculate and compute the performance of a regression model. It is the average of the squares of the errors, that are the differences in the actual and predicted values. Smaller the MSE, the better the model's predictions are, because it indicates less deviation from the actual values.

$$MSE = \frac{1}{n} \sum_{i=1}^n (\text{Actual}_i - \text{Predicted}_i)^2$$

MSE indicates how much is the deviation in model's predictions are from the actual values. If the MSE is low, the model's predictions are more aligned with the true values. If it's high, the model is less accurate.

Why Square the Errors?: Squaring the errors ensures that the model penalizes larger mistakes more heavily than smaller ones, which helps the model focus on minimizing large errors and improving overall accuracy.

2) r2_score : r2_score is used for regression problems. It is used when we want to know how well the Model fits the overall data and understand relationship between values of X and Y axis.

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_{\text{true},i} - y_{\text{pred},i})^2}{\sum_{i=1}^n (y_{\text{true},i} - \bar{y}_{\text{true}})^2}$$

Value of R² lies between 0 to 1 where 0 indicates bad relationship and values close to 1 indicates perfect fit. sklearn model helps to compute the values. For both linear regression and polynomial regression values which is higher is preferred.

For Polynomial regression, high r2_score value is good but watch MSE/RMSE carefully for the sign of overfitting, especially in higher polynomial degree.

In Conclusion R² provides a relative measure of how well the model captures the variance in the data. By using these two metrics, you can evaluate the performance of each model, compare different regression techniques, and understand whether the model is underfitting or overfitting the data. and Higher the R², more acceptable the model and lower the MSE better model.

In [60]:

```
# R squared error
score_1 = metrics.r2_score(Y_train, training_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_train, training_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)
```


```
R squared error : 0.9999969829984965
Mean Absolute Error : 0.011182523717974735
```

Visualizing the actual Prices and predicted prices

In [63]:

```
plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Predicted Price")
plt.show()
```

Actual Price vs Predicted Price



Now predicting on test data.

In [64]:

```
# accuracy for prediction on test data
test_data_prediction = model.predict(X_test)
```

In [65]:

```
# R squared error
score_1 = metrics.r2_score(Y_test, test_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_test, test_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)
```

```
R squared error : 0.9057837838492537
Mean Absolute Error : 1.8908873698290656
```

Conclusion

The Boston Housing dataset is used to predict house prices using, first, a baseline of linear regression and, later on, an improved model of XGBoost.

Linear Regression Model: To start with, we used a **Linear Regression** model. It gave us the following results: • **R-squared value: 0.8724** • **Mean Absolute Error (MAE): 2.7099** (on the test set)

These results show that the Linear Regression model generalizes fairly well to the task of house price prediction, with the model explaining about 87% of the variance in the data. The performance was far from perfect, though, with quite a significant room for improvement as evidenced by the relatively high MAE. That implied the model was not able to grasp the complex relationships between the features and the target variable, house prices.

XGBoost Model:

To overcome the limitations of Linear Regression, we have moved on to use the **XGBoost Regressor**-one of the power machine learning algorithms that can handle the non-linear relationship between target and predictor variables along with sophisticated patterns in the data. The result for XGBoost was as follows : Training R-squared value = 0.999997 Training MAE = 0.0112 Test R-squared value = 0.9058 Test MAE = 1.8909

The XGBoost model showed a massive improvement from Linear Regression, with the R-squared perfect in training. The performance on the test was also very strong, at an R-squared of 0.9058, indicating that over 90% of the variance in the test set was explained by the model. The MAE was way lower compared to Linear Regression, which indicated much more accurate predictions.

However, the XGBoost model's performance on the test set was not as perfect as the training set, which is typical of machine learning models. This slight decrease in performance suggests potential **overfitting**, where the model is overly tuned to the training data and may not generalize as well on unseen data.

Room for Improvement:

Despite the strong performance of the XGBoost model, there are still areas for improvement:

1. **Overfitting:** The model performed extremely well on the training set, whereas for the test set, performance was slightly lower, which might indicate overfitting. We can try techniques like **cross-validation** to help generalize better on unseen data or tune the hyperparameters.
2. **Hyperparameter Tuning:** The performance of XGBoost can be further enhanced by the use of different **hyperparameter optimization techniques** such as Grid Search or Random Search to find the most suitable set of parameters.
3. **Feature Engineering:** Inclusion of other relevant features or removal of a few irrelevant or redundant ones may further improve the accuracy. Feature interaction and transformation may also yield better results in improving the predictive capability of the model.
4. **Model Comparison:** Although XGBoost outperformed Linear Regression, it would be useful to compare its results with other models such as **Random Forests** or **Gradient Boosting Machines** just to verify that XGBoost is truly the best model for this problem.

Conclusion - We can see that we actually improved significantly in predicting house prices by moving from Linear Regression to XGBoost. The XGBoost model gave excellent performance, especially on lower error rates; however, steps for improving its generalization by preventing overfitting can include tuning its hyperparameters, applying regularization, and cross-validation. This project demonstrates the application of a more advanced machine learning model like XGBoost compared to the Linear Regression method; however, there can be a continuous refinement and optimization.

