

Introduction for the California Housing Price Prediction Project

In this project, we build and compare multiple machine learning models to predict housing prices in California. Using Python and libraries like Pandas, NumPy, Scikit-learn, and TensorFlow/Keras, to analyze California housing dataset. This dataset includes features such as geographical coordinates, housing attributes, and proximity to the ocean. The target variable is the median housing price in various districts.

The primary objective is to explore regression techniques to predict continuous values effectively. The project demonstrates the end-to-end machine learning workflow, including data preprocessing, feature scaling, model selection, and evaluation.

The goal of this project is to predict median house prices in California using various machine learning models. We experiment with traditional regression models and neural networks, comparing their performance to determine the best approach

IMPORTING NECESSARY LIBRARIES AND DEPENDENCIES

These libraries are essential for data handling (Pandas), numerical computations (NumPy), and data visualization (Matplotlib, Seaborn). Also We import machine learning models (LinearRegression, RandomForest, etc.), preprocessing tools (StandardScaler), and evaluation metrics (mean_squared_error) from Scikit-learn. While TensorFlow/Keras is used to design, train, and save neural networks.

In []:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.metrics import mean_squared_error as mse
from sklearn.neighbors import KNeighborsRegressor
```

Loading the Dataset

In []:

```
housing_pd = pd.read_csv('/content/housing.csv')
```

In []:

```
housing_pd
```

Out[]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_housing_price
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	
...
20635	-121.09	39.48	25.0	1665.0	374.0	845.0	330.0	1.5603	
20636	-121.21	39.49	18.0	697.0	150.0	356.0	114.0	2.5568	

20637	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median
20638	-121.32	39.43	18.0	1860.0	409.0	741.0	349.0	1.8672	
20639	-121.24	39.37	16.0	2785.0	616.0	1387.0	530.0	2.3886	

20640 rows × 10 columns



In []:

```
#To understand the distribution of the categorical variable ocean_proximity.
housing_pd['ocean_proximity'].value_counts()
```

Out[]:

count	
ocean_proximity	
<1H OCEAN	9136
INLAND	6551
NEAR OCEAN	2658
NEAR BAY	2290
ISLAND	5

dtype: int64

Shuffling and Encoding categorical variables

Shuffling ensures randomness in the dataset. Encoding converts ocean_proximity into numerical "dummy variables," which machine learning models can process.


In []:

```
housing_pd_shuffled = housing_pd.sample(n=len(housing_pd), random_state=1)
housing_pd_shuffled
```

Out[]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	media
4712	-118.36	34.06	39.0	2810.0	670.0	1109.0	624.0	3.2500	
2151	-119.78	36.78	37.0	2185.0	455.0	1143.0	438.0	1.9784	
15927	-122.42	37.73	46.0	1819.0	411.0	1534.0	406.0	4.0132	
82	-122.28	37.81	52.0	340.0	97.0	200.0	87.0	1.5208	
8161	-118.13	33.82	37.0	1530.0	290.0	711.0	283.0	5.1795	
...	
10955	-117.88	33.76	17.0	1768.0	474.0	1079.0	436.0	1.7823	
17289	-119.63	34.42	42.0	1765.0	263.0	753.0	260.0	8.5608	
5192	-118.26	33.93	42.0	1433.0	295.0	775.0	293.0	1.1326	
12172	-117.16	33.73	10.0	2381.0	454.0	1323.0	477.0	2.6322	
235	-122.20	37.79	35.0	1802.0	459.0	1009.0	390.0	2.3036	

20640 rows × 10 columns



In []:

```
# Preprocess ocean_proximity
housing_pd_shuffled = housing_pd.sample(n=len(housing_pd), random_state=1)
```

```
housing_pd_final = pd.concat([
    housing_pd_shuffled.drop('ocean_proximity', axis=1),
    pd.get_dummies(housing_pd_shuffled['ocean_proximity']),
    axis=1
])
```

In []:

```
# Rearrange columns
housing_pd_final = housing_pd_final[[
    'longitude', 'latitude', 'housing_median_age', 'total_rooms',
    'total_bedrooms', 'population', 'households', 'median_income',
    '<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN', 'median_house_value']]
```

In []:

```
housing_pd_final
```

Out[]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	<1H OCEAN
4712	-118.36	34.06	39.0	2810.0	670.0	1109.0	624.0	3.2500	True
2151	-119.78	36.78	37.0	2185.0	455.0	1143.0	438.0	1.9784	False
15927	-122.42	37.73	46.0	1819.0	411.0	1534.0	406.0	4.0132	False
82	-122.28	37.81	52.0	340.0	97.0	200.0	87.0	1.5208	False
8161	-118.13	33.82	37.0	1530.0	290.0	711.0	283.0	5.1795	True
...
10955	-117.88	33.76	17.0	1768.0	474.0	1079.0	436.0	1.7823	True
17289	-119.63	34.42	42.0	1765.0	263.0	753.0	260.0	8.5608	True
5192	-118.26	33.93	42.0	1433.0	295.0	775.0	293.0	1.1326	True
12172	-117.16	33.73	10.0	2381.0	454.0	1323.0	477.0	2.6322	False
235	-122.20	37.79	35.0	1802.0	459.0	1009.0	390.0	2.3036	False

20640 rows x 14 columns



Handeling Missing Values

In []:

```
housing_pd_final = housing_pd_final.dropna()
```

In []:

```
len(housing_pd_final)
```

Out[]:

20433

Splitting the data

Here Data is divided into:

- 1)Training set: Used to train the model.
- 2)Validation set: Used to fine-tune hyperparameters.
- 3)Test set: Used to evaluate the final model's performance.

In []:

```
train_pd, test_pd, val_pd = housing_pd_final[:18000], housing_pd_final[18000:19217], housing_pd_final[19217:]
len(train_pd), len(test_pd), len(val_pd)
```

```
Out[ ]:

(18000, 1217, 1218)
```

```
In [ ]:

X_train, y_train = train_pd.to_numpy()[:, :-1], train_pd.to_numpy()[:, -1]
X_val, y_val = val_pd.to_numpy()[:, :-1], val_pd.to_numpy()[:, -1]
X_test, y_test = test_pd.to_numpy()[:, :-1], test_pd.to_numpy()[:, -1]

X_train.shape, y_train.shape, X_val.shape, y_val.shape, X_test.shape, y_test.shape
```

```
Out[ ]:

((18000, 13), (18000,), (1218, 13), (1218,), (1217, 13), (1217,))
```

Feature Scaling & Custom Processor

Scaling ensures all numerical features have a similar range, preventing dominance of features with larger values.

while Preprocessor applies scaling consistently to training, validation, and test sets, ensuring uniform transformation.

```
In [ ]:

scaler = StandardScaler().fit(X_train[:, :8])

def preprocessor(X):
    A = np.copy(X)
    A[:, :8] = scaler.transform(A[:, :8])
    return A

X_train, X_val, X_test = preprocessor(X_train), preprocessor(X_val), preprocessor(X_test)
```

```
In [ ]:

X_train.shape, X_val.shape, X_test.shape
```

```
Out[ ]:

((18000, 13), (1218, 13), (1217, 13))
```

```
In [ ]:
```

```
In [ ]:

pd.DataFrame(X_train).head()
```

Out[]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.603443	-0.736073	0.820845	0.081039	0.315396	-0.27684	0.328234	-0.326667	True	False	False	False	False
1	-0.105122	0.537108	0.661774	-0.206526	-0.196843	-0.246809	-0.160526	-0.995001	False	True	False	False	False
2	-1.422454	0.981785	1.377592	-0.374924	-0.301674	0.098553	-0.244613	0.074459	False	False	False	True	False
3	-1.352596	1.019231	1.854804	-1.055419	-1.049782	-1.07974	-1.082862	-1.235508	False	False	False	True	False
4	0.718211	-0.848412	0.661774	-0.507894	-0.589957	-0.628385	-0.567825	0.687448	True	False	False	False	False


```
In [ ]:

pd.DataFrame(X_train).head()
```

```
pd.DataFrame(X_train)[0].hist()
```

Out[]:

<Axes: >



A)using Linear Regression

A simple model that predicts house prices by fitting a straight line through the data which May underperform with complex data or non-linear relationships.

In []:

```
lm = LinearRegression().fit(X_train, y_train)
mse(lm.predict(X_train), y_train, squared=False), mse(lm.predict(X_val), y_val, squared=False)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

Out[]:

(68593.05578127236, 71382.43558330165)

B)K-Nearest Neighbors(KNN)

Predicts prices based on the average values of the nearest 10 data points. Limitation: Can overfit if the number of neighbors is too small.

In []:

```
knn = KNeighborsRegressor(n_neighbors=10).fit(X_train, y_train)
mse(knn.predict(X_train), y_train, squared=False), mse(knn.predict(X_val), y_val, squared=False)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

```
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(
```

Out[]:

```
(53759.09908812057, 62161.22860469906)
```

C)Random Forest

Uses multiple decision trees to make predictions and Handles non-linear data well and reduces overfitting via ensemble learning.

In []:

```
rfr = RandomForestRegressor(max_depth=10).fit(X_train, y_train)
mse(rfr.predict(X_train), y_train, squared=False), mse(rfr.predict(X_val), y_val, squared=False)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

Out[]:

```
(43569.9006950773, 53183.737414903306)
```

D)Gradient Boosting

Combines weak learners incrementally to improve accuracy. Typically outperforms Random Forest in predictive tasks.

In []:

```
gbr = GradientBoostingRegressor(n_estimators=250).fit(X_train, y_train)
mse(gbr.predict(X_train), y_train, squared=False), mse(gbr.predict(X_val), y_val, squared=False)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

Out[]:

```
(47274.82259072158, 51210.606733472116)
```

In []:

```
# Convert data to float32
X_train = X_train.astype('float32')
X_val = X_val.astype('float32')
y_train = y_train.astype('float32')
y_val = y_val.astype('float32')
```

IMPORTING NECESSARY LIBRARIES AND DEPENDENCIES.

1)for Building Simple Neural Network**

A basic neural network with one hidden layer and two neurons.


```
ln | ] :
```

```
from tensorflow.keras.models import Sequential,load_model
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.metrics import RootMeanSquaredError
from tensorflow.keras.optimizers import Adam


simple_nn = Sequential()
simple_nn.add(InputLayer((13,)))
simple_nn.add(Dense(2, 'relu'))
simple_nn.add(Dense(1, 'linear'))

opt = Adam(learning_rate=.1)
cp = ModelCheckpoint('models/simple_nn.keras', save_best_only=True) #save as .keras form at
simple_nn.compile(optimizer=opt, loss='mse', metrics=[RootMeanSquaredError()])
simple_nn.fit(x=X_train, y=y_train, validation_data=(X_val, y_val), callbacks=[cp], epochs=100)
```


Epoch 1/100

563/563  2s 2ms/step - loss: 54539640832.0000 - root_mean_squared_error: 233498.9531 - val_loss: 41773555712.0000 - val_root_mean_squared_error: 204385.7969


Epoch 2/100

563/563  1s 2ms/step - loss: 34940604416.0000 - root_mean_squared_error: 186656.0312 - val_loss: 18346248192.0000 - val_root_mean_squared_error: 135448.3281


Epoch 3/100

563/563  1s 2ms/step - loss: 14725444608.0000 - root_mean_squared_error: 121149.6953 - val_loss: 9118205952.0000 - val_root_mean_squared_error: 95489.2969


Epoch 4/100

563/563  1s 2ms/step - loss: 8249291776.0000 - root_mean_squared_error: 90765.3516 - val_loss: 7016109568.0000 - val_root_mean_squared_error: 83762.2188


Epoch 5/100

563/563  1s 2ms/step - loss: 6735520256.0000 - root_mean_squared_error: 82020.9297 - val_loss: 6083796480.0000 - val_root_mean_squared_error: 77998.6953


Epoch 6/100

563/563  1s 2ms/step - loss: 5585117696.0000 - root_mean_squared_error: 74721.6016 - val_loss: 5544328704.0000 - val_root_mean_squared_error: 74460.2500


Epoch 7/100

563/563  2s 3ms/step - loss: 5204871168.0000 - root_mean_squared_error: 72125.4219 - val_loss: 5293306368.0000 - val_root_mean_squared_error: 72755.1094


Epoch 8/100

563/563  2s 2ms/step - loss: 4790865408.0000 - root_mean_squared_error: 69204.6719 - val_loss: 5179879424.0000 - val_root_mean_squared_error: 71971.3828


Epoch 9/100

563/563  1s 2ms/step - loss: 4920813056.0000 - root_mean_squared_error: 70141.4219 - val_loss: 5137507840.0000 - val_root_mean_squared_error: 71676.4141


Epoch 10/100

563/563  1s 2ms/step - loss: 4634716160.0000 - root_mean_squared_error: 68049.6016 - val_loss: 5105841152.0000 - val_root_mean_squared_error: 71455.1719


Epoch 11/100

563/563  1s 2ms/step - loss: 4645094400.0000 - root_mean_squared_error: 68147.4766 - val_loss: 5082308608.0000 - val_root_mean_squared_error: 71290.3125


Epoch 12/100

563/563  1s 2ms/step - loss: 4717749248.0000 - root_mean_squared_error: 68681.7344 - val_loss: 5066271744.0000 - val_root_mean_squared_error: 71177.7500


Epoch 13/100

563/563  1s 2ms/step - loss: 4613076992.0000 - root_mean_squared_error: 67911.8516 - val_loss: 5059229184.0000 - val_root_mean_squared_error: 71128.2578


Epoch 14/100

563/563  1s 2ms/step - loss: 4455755776.0000 - root_mean_squared_error: 66721.1484 - val_loss: 5058302464.0000 - val_root_mean_squared_error: 71121.7422


Epoch 15/100

563/563  1s 2ms/step - loss: 4638859776.0000 - root_mean_squared_error: 68104.7422 - val_loss: 5044654592.0000 - val_root_mean_squared_error: 71025.7344

Epoch 16/100










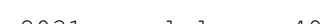

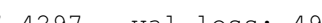
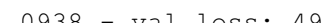


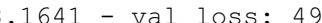

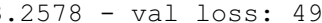
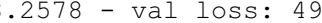
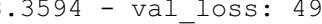
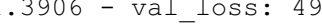

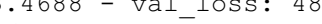

563/563  2s 2ms/step - loss: 4555708928.0000 - root_mean_squared_error: 67487.5391 - val_loss: 5044612096.0000 - val_root_mean_squared_error: 71025.4297

Epoch 17/100

563/563  1s 2ms/step - loss: 4690583552.0000 - root_mean_squared_error: 68484.6562 - val_loss: 5053232640.0000 - val_root_mean_squared_error: 71086.0938

Epoch 18/100

563/563  2s 2ms/step - loss: 4631273472.0000 - root mean squared error

r: 68047.0234 - val_loss: 5029160448.0000 - val_root_mean_squared_error: 70916.5703
Epoch 19/100
563/563  1s 2ms/step - loss: 4499521536.0000 - root_mean_squared_error: 67056.4922 - val_loss: 5035070976.0000 - val_root_mean_squared_error: 70958.2344
Epoch 20/100
563/563  1s 2ms/step - loss: 4797551616.0000 - root_mean_squared_error: 69251.0391 - val_loss: 5032795648.0000 - val_root_mean_squared_error: 70942.2031
Epoch 21/100
563/563  1s 2ms/step - loss: 4562572288.0000 - root_mean_squared_error: 67525.8828 - val_loss: 5033680896.0000 - val_root_mean_squared_error: 70948.4375
Epoch 22/100
563/563  1s 2ms/step - loss: 4501138944.0000 - root_mean_squared_error: 67079.5312 - val_loss: 5014000128.0000 - val_root_mean_squared_error: 70809.6016
Epoch 23/100
563/563  1s 2ms/step - loss: 4566908416.0000 - root_mean_squared_error: 67565.7188 - val_loss: 5000335872.0000 - val_root_mean_squared_error: 70713.0547
Epoch 24/100
563/563  1s 2ms/step - loss: 4560048640.0000 - root_mean_squared_error: 67522.9766 - val_loss: 4997998080.0000 - val_root_mean_squared_error: 70696.5234
Epoch 25/100
563/563  1s 2ms/step - loss: 4510138368.0000 - root_mean_squared_error: 67145.5312 - val_loss: 5007716864.0000 - val_root_mean_squared_error: 70765.2266
Epoch 26/100
563/563  2s 2ms/step - loss: 4602163712.0000 - root_mean_squared_error: 67835.9141 - val_loss: 5010356224.0000 - val_root_mean_squared_error: 70783.8672
Epoch 27/100
563/563  1s 2ms/step - loss: 4547445248.0000 - root_mean_squared_error: 67426.8125 - val_loss: 4978567680.0000 - val_root_mean_squared_error: 70558.9688
Epoch 28/100
563/563  1s 2ms/step - loss: 4559151104.0000 - root_mean_squared_error: 67515.2031 - val_loss: 4975286784.0000 - val_root_mean_squared_error: 70535.7109
Epoch 29/100
563/563  1s 2ms/step - loss: 4496536576.0000 - root_mean_squared_error: 67046.3047 - val_loss: 4982700032.0000 - val_root_mean_squared_error: 70588.2422
Epoch 30/100
563/563  1s 2ms/step - loss: 4587983872.0000 - root_mean_squared_error: 67727.4297 - val_loss: 4969357312.0000 - val_root_mean_squared_error: 70493.6719
Epoch 31/100
563/563  1s 2ms/step - loss: 4791488000.0000 - root_mean_squared_error: 69181.0938 - val_loss: 4949418496.0000 - val_root_mean_squared_error: 70352.1016
Epoch 32/100
563/563  1s 2ms/step - loss: 4594275328.0000 - root_mean_squared_error: 67778.4844 - val_loss: 4962390528.0000 - val_root_mean_squared_error: 70444.2344
Epoch 33/100
563/563  1s 2ms/step - loss: 4399966208.0000 - root_mean_squared_error: 66321.8594 - val_loss: 4965132800.0000 - val_root_mean_squared_error: 70463.6953
Epoch 34/100
563/563  1s 2ms/step - loss: 4614180864.0000 - root_mean_squared_error: 67918.1641 - val_loss: 4946462208.0000 - val_root_mean_squared_error: 70331.0938
Epoch 35/100
563/563  1s 2ms/step - loss: 4390273536.0000 - root_mean_squared_error: 66254.6250 - val_loss: 4944105984.0000 - val_root_mean_squared_error: 70314.3359
Epoch 36/100
563/563  2s 2ms/step - loss: 4382669312.0000 - root_mean_squared_error: 66193.2578 - val_loss: 4909303296.0000 - val_root_mean_squared_error: 70066.4219
Epoch 37/100
563/563  1s 3ms/step - loss: 4607254016.0000 - root_mean_squared_error: 67868.2578 - val_loss: 4921146880.0000 - val_root_mean_squared_error: 70150.8828
Epoch 38/100
563/563  1s 2ms/step - loss: 4531285504.0000 - root_mean_squared_error: 67308.3594 - val_loss: 4902928384.0000 - val_root_mean_squared_error: 70020.9141
Epoch 39/100
563/563  2s 1ms/step - loss: 4444804096.0000 - root_mean_squared_error: 66641.3906 - val_loss: 4904748544.0000 - val_root_mean_squared_error: 70033.9062
Epoch 40/100
563/563  1s 2ms/step - loss: 4475655168.0000 - root_mean_squared_error: 66892.1406 - val_loss: 4873295360.0000 - val_root_mean_squared_error: 69808.9922
Epoch 41/100
563/563  1s 2ms/step - loss: 4526555648.0000 - root_mean_squared_error: 67273.4688 - val_loss: 4886693376.0000 - val_root_mean_squared_error: 69904.8906
Epoch 42/100
563/563  1s 2ms/step - loss: 4438777344.0000 - root mean squared error: 67046.3047 - val_loss: 4982700032.0000 - val_root_mean_squared_error: 70588.2422

r: 66620.0234 - val_loss: 4863424000.0000 - val_root_mean_squared_error: 69738.2500
Epoch 43/100
563/563 1s 2ms/step - loss: 4529377280.0000 - root_mean_squared_error: 67296.2734 - val_loss: 4862021120.0000 - val_root_mean_squared_error: 69728.1953
Epoch 44/100
563/563 1s 2ms/step - loss: 4472200192.0000 - root_mean_squared_error: 66858.0000 - val_loss: 4856575488.0000 - val_root_mean_squared_error: 69689.1328
Epoch 45/100
563/563 1s 2ms/step - loss: 4572540416.0000 - root_mean_squared_error: 67595.7578 - val_loss: 4859686400.0000 - val_root_mean_squared_error: 69711.4531
Epoch 46/100
563/563 1s 2ms/step - loss: 4486858752.0000 - root_mean_squared_error: 66978.5234 - val_loss: 4835853312.0000 - val_root_mean_squared_error: 69540.2969
Epoch 47/100
563/563 2s 2ms/step - loss: 4441146368.0000 - root_mean_squared_error: 66638.9844 - val_loss: 4842146304.0000 - val_root_mean_squared_error: 69585.5312
Epoch 48/100
563/563 2s 2ms/step - loss: 4439943680.0000 - root_mean_squared_error: 66622.0234 - val_loss: 4845543936.0000 - val_root_mean_squared_error: 69609.9375
Epoch 49/100
563/563 1s 2ms/step - loss: 4431044608.0000 - root_mean_squared_error: 66549.9219 - val_loss: 4802679296.0000 - val_root_mean_squared_error: 69301.3672
Epoch 50/100
563/563 1s 2ms/step - loss: 4457390592.0000 - root_mean_squared_error: 66756.8906 - val_loss: 4829318656.0000 - val_root_mean_squared_error: 69493.2969
Epoch 51/100
563/563 1s 2ms/step - loss: 4307536896.0000 - root_mean_squared_error: 65626.4453 - val_loss: 4807420416.0000 - val_root_mean_squared_error: 69335.5625
Epoch 52/100
563/563 1s 2ms/step - loss: 4259789056.0000 - root_mean_squared_error: 65260.5898 - val_loss: 4793296384.0000 - val_root_mean_squared_error: 69233.6328
Epoch 53/100
563/563 1s 2ms/step - loss: 4422423040.0000 - root_mean_squared_error: 66490.9844 - val_loss: 4811170816.0000 - val_root_mean_squared_error: 69362.6016
Epoch 54/100
563/563 1s 2ms/step - loss: 4500658688.0000 - root_mean_squared_error: 67070.8438 - val_loss: 4789377536.0000 - val_root_mean_squared_error: 69205.3281
Epoch 55/100
563/563 1s 2ms/step - loss: 4304159232.0000 - root_mean_squared_error: 65601.5312 - val_loss: 4777832448.0000 - val_root_mean_squared_error: 69121.8672
Epoch 56/100
563/563 2s 3ms/step - loss: 4398196224.0000 - root_mean_squared_error: 66304.4375 - val_loss: 4777438720.0000 - val_root_mean_squared_error: 69119.0156
Epoch 57/100
563/563 3s 3ms/step - loss: 4544287232.0000 - root_mean_squared_error: 67397.7500 - val_loss: 4761050624.0000 - val_root_mean_squared_error: 69000.3672
Epoch 58/100
563/563 2s 2ms/step - loss: 4450718208.0000 - root_mean_squared_error: 66707.0781 - val_loss: 4784960000.0000 - val_root_mean_squared_error: 69173.4062
Epoch 59/100
563/563 1s 2ms/step - loss: 4431842304.0000 - root_mean_squared_error: 66542.6797 - val_loss: 4785341440.0000 - val_root_mean_squared_error: 69176.1641
Epoch 60/100
563/563 1s 2ms/step - loss: 4287498752.0000 - root_mean_squared_error: 65451.5156 - val_loss: 4753483264.0000 - val_root_mean_squared_error: 68945.5078
Epoch 61/100
563/563 1s 2ms/step - loss: 4450786816.0000 - root_mean_squared_error: 66710.0938 - val_loss: 4743922688.0000 - val_root_mean_squared_error: 68876.1406
Epoch 62/100
563/563 1s 2ms/step - loss: 4459515904.0000 - root_mean_squared_error: 66775.8594 - val_loss: 4735197696.0000 - val_root_mean_squared_error: 68812.7734
Epoch 63/100
563/563 1s 2ms/step - loss: 4416309760.0000 - root_mean_squared_error: 66443.1641 - val_loss: 4752134144.0000 - val_root_mean_squared_error: 68935.7266
Epoch 64/100
563/563 1s 2ms/step - loss: 4311967744.0000 - root_mean_squared_error: 65660.7891 - val_loss: 4742831104.0000 - val_root_mean_squared_error: 68868.2188
Epoch 65/100
563/563 1s 2ms/step - loss: 4292076800.0000 - root_mean_squared_error: 65511.1094 - val_loss: 4729505792.0000 - val_root_mean_squared_error: 68771.4062
Epoch 66/100
563/563 2s 2ms/step - loss: 4455957504.0000 - root mean squared error

r: 66730.6172 - val_loss: 4739006976.0000 - val_root_mean_squared_error: 68840.4453
Epoch 67/100
563/563 1s 3ms/step - loss: 4397680128.0000 - root_mean_squared_error: 66299.3281 - val_loss: 4713782784.0000 - val_root_mean_squared_error: 68656.9922
Epoch 68/100
563/563 1s 2ms/step - loss: 4305144832.0000 - root_mean_squared_error: 65596.6016 - val_loss: 4700944896.0000 - val_root_mean_squared_error: 68563.4375
Epoch 69/100
563/563 2s 2ms/step - loss: 4272348672.0000 - root_mean_squared_error: 65358.3047 - val_loss: 4702661120.0000 - val_root_mean_squared_error: 68575.9531
Epoch 70/100
563/563 1s 2ms/step - loss: 4496883712.0000 - root_mean_squared_error: 67040.0156 - val_loss: 4701617152.0000 - val_root_mean_squared_error: 68568.3359
Epoch 71/100
563/563 1s 2ms/step - loss: 4281522432.0000 - root_mean_squared_error: 65429.9180 - val_loss: 4705542656.0000 - val_root_mean_squared_error: 68596.9609
Epoch 72/100
563/563 1s 2ms/step - loss: 4392652288.0000 - root_mean_squared_error: 66248.1250 - val_loss: 4678725632.0000 - val_root_mean_squared_error: 68401.2109
Epoch 73/100
563/563 1s 2ms/step - loss: 4228637440.0000 - root_mean_squared_error: 65016.9023 - val_loss: 4679358464.0000 - val_root_mean_squared_error: 68405.8359
Epoch 74/100
563/563 1s 2ms/step - loss: 4343191040.0000 - root_mean_squared_error: 65896.5000 - val_loss: 4663543808.0000 - val_root_mean_squared_error: 68290.1406
Epoch 75/100
563/563 1s 2ms/step - loss: 4283275008.0000 - root_mean_squared_error: 65443.3281 - val_loss: 4677759488.0000 - val_root_mean_squared_error: 68394.1484
Epoch 76/100
563/563 2s 2ms/step - loss: 4281000448.0000 - root_mean_squared_error: 65425.3320 - val_loss: 4647067136.0000 - val_root_mean_squared_error: 68169.3984
Epoch 77/100
563/563 2s 2ms/step - loss: 4335552512.0000 - root_mean_squared_error: 65823.2734 - val_loss: 4649236480.0000 - val_root_mean_squared_error: 68185.3125
Epoch 78/100
563/563 1s 2ms/step - loss: 4189521152.0000 - root_mean_squared_error: 64711.1641 - val_loss: 4629218816.0000 - val_root_mean_squared_error: 68038.3594
Epoch 79/100
563/563 1s 2ms/step - loss: 4243095040.0000 - root_mean_squared_error: 65132.8594 - val_loss: 4641016320.0000 - val_root_mean_squared_error: 68125.0078
Epoch 80/100
563/563 1s 2ms/step - loss: 4345121792.0000 - root_mean_squared_error: 65905.9609 - val_loss: 4659392512.0000 - val_root_mean_squared_error: 68259.7422
Epoch 81/100
563/563 1s 2ms/step - loss: 4396019200.0000 - root_mean_squared_error: 66294.7422 - val_loss: 4643544064.0000 - val_root_mean_squared_error: 68143.5547
Epoch 82/100
563/563 1s 2ms/step - loss: 4285911040.0000 - root_mean_squared_error: 65454.3555 - val_loss: 4642816000.0000 - val_root_mean_squared_error: 68138.2109
Epoch 83/100
563/563 1s 2ms/step - loss: 4263274496.0000 - root_mean_squared_error: 65276.7070 - val_loss: 4611879424.0000 - val_root_mean_squared_error: 67910.8203
Epoch 84/100
563/563 1s 2ms/step - loss: 4195711232.0000 - root_mean_squared_error: 64763.8086 - val_loss: 4605697024.0000 - val_root_mean_squared_error: 67865.2891
Epoch 85/100
563/563 2s 2ms/step - loss: 4130084864.0000 - root_mean_squared_error: 64248.7773 - val_loss: 4629566976.0000 - val_root_mean_squared_error: 68040.9219
Epoch 86/100
563/563 2s 3ms/step - loss: 4174304256.0000 - root_mean_squared_error: 64597.7461 - val_loss: 4603520000.0000 - val_root_mean_squared_error: 67849.2422
Epoch 87/100
563/563 2s 2ms/step - loss: 4204616704.0000 - root_mean_squared_error: 64834.9023 - val_loss: 4597460992.0000 - val_root_mean_squared_error: 67804.5781
Epoch 88/100
563/563 1s 2ms/step - loss: 4314948608.0000 - root_mean_squared_error: 65679.0078 - val_loss: 4628914688.0000 - val_root_mean_squared_error: 68036.1250
Epoch 89/100
563/563 1s 2ms/step - loss: 4340725248.0000 - root_mean_squared_error: 65879.4609 - val_loss: 4597753856.0000 - val_root_mean_squared_error: 67806.7422
Epoch 90/100
563/563 1s 2ms/step - loss: 4320001024.0000 - root mean squared error

```

r: 65714.3438 - val_loss: 4596164608.0000 - val_root_mean_squared_error: 67795.0156
Epoch 91/100
563/563 ————— 1s 2ms/step - loss: 4308797952.0000 - root_mean_squared_error: 65628.4062 - val_loss: 4600026112.0000 - val_root_mean_squared_error: 67823.4922
Epoch 92/100
563/563 ————— 1s 2ms/step - loss: 4294361600.0000 - root_mean_squared_error: 65529.2266 - val_loss: 4600787968.0000 - val_root_mean_squared_error: 67829.1094
Epoch 93/100
563/563 ————— 1s 2ms/step - loss: 4311780352.0000 - root_mean_squared_error: 65659.9688 - val_loss: 4594822656.0000 - val_root_mean_squared_error: 67785.1250
Epoch 94/100
563/563 ————— 1s 2ms/step - loss: 4289721856.0000 - root_mean_squared_error: 65487.6602 - val_loss: 4602252800.0000 - val_root_mean_squared_error: 67839.9062
Epoch 95/100
563/563 ————— 2s 3ms/step - loss: 4193928960.0000 - root_mean_squared_error: 64746.9375 - val_loss: 4600219136.0000 - val_root_mean_squared_error: 67824.9141
Epoch 96/100
563/563 ————— 2s 2ms/step - loss: 4257855232.0000 - root_mean_squared_error: 65236.8750 - val_loss: 4589726208.0000 - val_root_mean_squared_error: 67747.5156
Epoch 97/100
563/563 ————— 1s 2ms/step - loss: 4367850496.0000 - root_mean_squared_error: 66035.8672 - val_loss: 4570818048.0000 - val_root_mean_squared_error: 67607.8281
Epoch 98/100
563/563 ————— 1s 2ms/step - loss: 4291658752.0000 - root_mean_squared_error: 65482.3906 - val_loss: 4583065088.0000 - val_root_mean_squared_error: 67698.3359
Epoch 99/100
563/563 ————— 1s 2ms/step - loss: 4103278848.0000 - root_mean_squared_error: 64045.1992 - val_loss: 4571387904.0000 - val_root_mean_squared_error: 67612.0391
Epoch 100/100
563/563 ————— 1s 2ms/step - loss: 4221046016.0000 - root_mean_squared_error: 64949.8672 - val_loss: 4567960576.0000 - val_root_mean_squared_error: 67586.6875

```

Out[]:

<keras.src.callbacks.history.History at 0x7880582c0df0>

In []:

```

simple_nn = load_model('models/simple_nn.keras')
mse(simple_nn.predict(X_train), y_train, squared=False), mse(simple_nn.predict(X_val), y_val, squared=False)

```

```

563/563 ————— 1s 1ms/step
39/39 ————— 0s 1ms/step

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(

```

Out[]:

(65134.406, 67586.695)

2)for Medium Neural Network

Adds more complexity with two hidden layers (32 and 16 neurons)

In []:

```

medium_nn = Sequential()
medium_nn.add(InputLayer((13,)))
medium_nn.add(Dense(32, 'relu'))
medium_nn.add(Dense(16, 'relu'))
medium_nn.add(Dense(1, 'linear'))

opt = Adam(learning_rate=.1)
cp = ModelCheckpoint('models/medium_nn.keras', save_best_only=True)

```

```
medium_nn.compile(optimizer=opt, loss='mse', metrics=[RootMeanSquaredError()])
medium_nn.fit(x=X_train, y=y_train, validation_data=(X_val, y_val), callbacks=[cp], epoch=100)
```

Epoch 1/100

563/563 ————— 3s 3ms/step - loss: 19243558912.0000 - root_mean_squared_error: 132632.7500 - val_loss: 5051724288.0000 - val_root_mean_squared_error: 71075.4844

Epoch 2/100

563/563 ————— 2s 3ms/step - loss: 4470901760.0000 - root_mean_squared_error: 66862.2266 - val_loss: 4687271424.0000 - val_root_mean_squared_error: 68463.6484

Epoch 3/100

563/563 ————— 2s 2ms/step - loss: 4235913472.0000 - root_mean_squared_error: 65050.3984 - val_loss: 4779671552.0000 - val_root_mean_squared_error: 69135.1719

Epoch 4/100

563/563 ————— 3s 4ms/step - loss: 4214558720.0000 - root_mean_squared_error: 64906.3906 - val_loss: 4535901696.0000 - val_root_mean_squared_error: 67349.1016

Epoch 5/100

563/563 ————— 2s 2ms/step - loss: 4208945664.0000 - root_mean_squared_error: 64865.4570 - val_loss: 4441363968.0000 - val_root_mean_squared_error: 66643.5625

Epoch 6/100

563/563 ————— 3s 3ms/step - loss: 4287689472.0000 - root_mean_squared_error: 65476.4062 - val_loss: 4505629696.0000 - val_root_mean_squared_error: 67123.9844

Epoch 7/100

563/563 ————— 2s 2ms/step - loss: 4231639808.0000 - root_mean_squared_error: 65039.0117 - val_loss: 4413633536.0000 - val_root_mean_squared_error: 66435.1797

Epoch 8/100

563/563 ————— 2s 3ms/step - loss: 4149835520.0000 - root_mean_squared_error: 64411.9062 - val_loss: 4415021056.0000 - val_root_mean_squared_error: 66445.6250

Epoch 9/100

563/563 ————— 1s 3ms/step - loss: 4095428352.0000 - root_mean_squared_error: 63988.2031 - val_loss: 4243283968.0000 - val_root_mean_squared_error: 65140.4922

Epoch 10/100

563/563 ————— 2s 2ms/step - loss: 4011297536.0000 - root_mean_squared_error: 63327.5742 - val_loss: 4294592256.0000 - val_root_mean_squared_error: 65533.1367

Epoch 11/100

563/563 ————— 1s 2ms/step - loss: 4146739200.0000 - root_mean_squared_error: 64383.9141 - val_loss: 4316926464.0000 - val_root_mean_squared_error: 65703.3203

Epoch 12/100

563/563 ————— 1s 2ms/step - loss: 3931311360.0000 - root_mean_squared_error: 62678.6016 - val_loss: 4340554752.0000 - val_root_mean_squared_error: 65882.8828

Epoch 13/100

563/563 ————— 1s 2ms/step - loss: 4063308800.0000 - root_mean_squared_error: 63733.9688 - val_loss: 4332282368.0000 - val_root_mean_squared_error: 65820.0781

Epoch 14/100

563/563 ————— 1s 2ms/step - loss: 3925444096.0000 - root_mean_squared_error: 62646.2617 - val_loss: 4260986624.0000 - val_root_mean_squared_error: 65276.2344

Epoch 15/100

563/563 ————— 1s 2ms/step - loss: 3991703040.0000 - root_mean_squared_error: 63172.4375 - val_loss: 4311772672.0000 - val_root_mean_squared_error: 65664.0859

Epoch 16/100

563/563 ————— 1s 2ms/step - loss: 4127987712.0000 - root_mean_squared_error: 64226.8750 - val_loss: 4260505600.0000 - val_root_mean_squared_error: 65272.5469

Epoch 17/100

563/563 ————— 2s 3ms/step - loss: 3937495808.0000 - root_mean_squared_error: 62745.0742 - val_loss: 4289468416.0000 - val_root_mean_squared_error: 65494.0352

Epoch 18/100

563/563 ————— 2s 2ms/step - loss: 3990547712.0000 - root_mean_squared_error: 63160.7188 - val_loss: 4455271936.0000 - val_root_mean_squared_error: 66747.8203

Epoch 19/100

563/563 ————— 1s 2ms/step - loss: 4035778816.0000 - root_mean_squared_error: 63519.3789 - val_loss: 4305367040.0000 - val_root_mean_squared_error: 65615.2969

Epoch 20/100

563/563 ————— 1s 2ms/step - loss: 4104944896.0000 - root_mean_squared_error: 64061.6406 - val_loss: 4284664064.0000 - val_root_mean_squared_error: 65457.3438

Epoch 21/100



















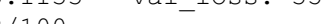
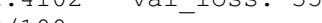
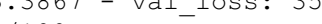
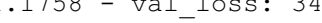
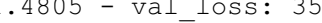
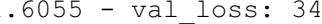
563/563 ————— 1s 2ms/step - loss: 3989139968.0000 - root_mean_squared_error: 63148.9648 - val_loss: 4303666688.0000 - val_root_mean_squared_error: 65602.3359

























Epoch 22/100

























563/563 ————— 1s 2ms/step - loss: 4117455104.0000 - root_mean_squared_error: 64154.7969 - val_loss: 4634318336.0000 - val_root_mean_squared_error: 68075.8281

Epoch 23/100

563/563 ————— 1s 2ms/step - loss: 3940097536.0000 - root_mean_squared_error: 63400.9753 - val_loss: 430097536.0000 - val_root_mean_squared_error: 65615.2969

r: 62767.7305 - val_loss: 4259236608.0000 - val_root_mean_squared_error: 65262.8281
Epoch 24/100
563/563  1s 2ms/step - loss: 3827083264.0000 - root_mean_squared_error: 61811.9609 - val_loss: 4191153152.0000 - val_root_mean_squared_error: 64739.1172
Epoch 25/100
563/563  1s 2ms/step - loss: 4119893760.0000 - root_mean_squared_error: 64167.0195 - val_loss: 4347846144.0000 - val_root_mean_squared_error: 65938.2031
Epoch 26/100
563/563  1s 2ms/step - loss: 3881918464.0000 - root_mean_squared_error: 62282.4531 - val_loss: 4245288192.0000 - val_root_mean_squared_error: 65155.8750
Epoch 27/100
563/563  2s 3ms/step - loss: 3891435264.0000 - root_mean_squared_error: 62360.9375 - val_loss: 4392724992.0000 - val_root_mean_squared_error: 66277.6328
Epoch 28/100
563/563  2s 2ms/step - loss: 3850518016.0000 - root_mean_squared_error: 62037.5703 - val_loss: 4091012352.0000 - val_root_mean_squared_error: 63961.0234
Epoch 29/100
563/563  1s 2ms/step - loss: 3861900288.0000 - root_mean_squared_error: 62134.2930 - val_loss: 4082724864.0000 - val_root_mean_squared_error: 63896.2031
Epoch 30/100
563/563  1s 2ms/step - loss: 3839528448.0000 - root_mean_squared_error: 61948.6484 - val_loss: 4060268032.0000 - val_root_mean_squared_error: 63720.2344
Epoch 31/100
563/563  1s 2ms/step - loss: 3783179008.0000 - root_mean_squared_error: 61499.1680 - val_loss: 3976538624.0000 - val_root_mean_squared_error: 63059.8008
Epoch 32/100
563/563  1s 2ms/step - loss: 3713291264.0000 - root_mean_squared_error: 60918.9844 - val_loss: 4084550144.0000 - val_root_mean_squared_error: 63910.4844
Epoch 33/100
563/563  1s 2ms/step - loss: 3541713664.0000 - root_mean_squared_error: 59499.5117 - val_loss: 3814196992.0000 - val_root_mean_squared_error: 61759.1836
Epoch 34/100
563/563  1s 2ms/step - loss: 3473965568.0000 - root_mean_squared_error: 58936.3125 - val_loss: 3815220480.0000 - val_root_mean_squared_error: 61767.4727
Epoch 35/100
563/563  1s 2ms/step - loss: 3447288576.0000 - root_mean_squared_error: 58709.3945 - val_loss: 3837622016.0000 - val_root_mean_squared_error: 61948.5430
Epoch 36/100
563/563  2s 3ms/step - loss: 3396995584.0000 - root_mean_squared_error: 58267.9570 - val_loss: 3692597504.0000 - val_root_mean_squared_error: 60766.7461
Epoch 37/100
563/563  2s 3ms/step - loss: 3329153024.0000 - root_mean_squared_error: 57675.5938 - val_loss: 3614981376.0000 - val_root_mean_squared_error: 60124.7148
Epoch 38/100
563/563  2s 2ms/step - loss: 3297738496.0000 - root_mean_squared_error: 57417.2227 - val_loss: 3780383744.0000 - val_root_mean_squared_error: 61484.8242
Epoch 39/100
563/563  1s 2ms/step - loss: 3337052672.0000 - root_mean_squared_error: 57760.2891 - val_loss: 3676929792.0000 - val_root_mean_squared_error: 60637.6914
Epoch 40/100
563/563  1s 2ms/step - loss: 3301145600.0000 - root_mean_squared_error: 57443.7188 - val_loss: 3597488384.0000 - val_root_mean_squared_error: 59979.0664
Epoch 41/100
563/563  1s 2ms/step - loss: 3162695424.0000 - root_mean_squared_error: 56233.1133 - val_loss: 3566669824.0000 - val_root_mean_squared_error: 59721.6016
Epoch 42/100
563/563  1s 2ms/step - loss: 3177044992.0000 - root_mean_squared_error: 56362.4102 - val_loss: 3536328704.0000 - val_root_mean_squared_error: 59467.0391
Epoch 43/100
563/563  1s 2ms/step - loss: 3268365824.0000 - root_mean_squared_error: 57143.3867 - val_loss: 3583668992.0000 - val_root_mean_squared_error: 59863.7539
Epoch 44/100
563/563  1s 2ms/step - loss: 3094115840.0000 - root_mean_squared_error: 55611.1758 - val_loss: 3460246528.0000 - val_root_mean_squared_error: 58823.8594
Epoch 45/100
563/563  1s 2ms/step - loss: 3105144576.0000 - root_mean_squared_error: 55711.4805 - val_loss: 3523337216.0000 - val_root_mean_squared_error: 59357.7070
Epoch 46/100
563/563  2s 3ms/step - loss: 3067174144.0000 - root_mean_squared_error: 55361.6055 - val_loss: 3438238720.0000 - val_root_mean_squared_error: 58636.4961
Epoch 47/100
563/563  2s 2ms/step - loss: 3078836224.0000 - root_mean_squared_error: 55361.6055 - val_loss: 3438238720.0000 - val_root_mean_squared_error: 58636.4961

r: 55476.2773 - val_loss: 3628026368.0000 - val_root_mean_squared_error: 60233.1016
Epoch 48/100
563/563  1s 2ms/step - loss: 3000899584.0000 - root_mean_squared_error: 54762.2031 - val_loss: 3434891264.0000 - val_root_mean_squared_error: 58607.9453
Epoch 49/100
563/563  1s 2ms/step - loss: 3115786496.0000 - root_mean_squared_error: 55798.1094 - val_loss: 3468494592.0000 - val_root_mean_squared_error: 58893.9258
Epoch 50/100
563/563  1s 2ms/step - loss: 3023492864.0000 - root_mean_squared_error: 54976.3477 - val_loss: 3419905536.0000 - val_root_mean_squared_error: 58479.9570
Epoch 51/100
563/563  1s 2ms/step - loss: 3084845056.0000 - root_mean_squared_error: 55530.3164 - val_loss: 3402247680.0000 - val_root_mean_squared_error: 58328.7891
Epoch 52/100
563/563  1s 2ms/step - loss: 3087880448.0000 - root_mean_squared_error: 55561.1797 - val_loss: 3348767744.0000 - val_root_mean_squared_error: 57868.5391
Epoch 53/100
563/563  1s 2ms/step - loss: 2907374592.0000 - root_mean_squared_error: 53887.1172 - val_loss: 3699054080.0000 - val_root_mean_squared_error: 60819.8477
Epoch 54/100
563/563  1s 2ms/step - loss: 3145741312.0000 - root_mean_squared_error: 56050.2734 - val_loss: 3385556992.0000 - val_root_mean_squared_error: 58185.5391
Epoch 55/100
563/563  1s 2ms/step - loss: 2900940544.0000 - root_mean_squared_error: 53845.1445 - val_loss: 3328689152.0000 - val_root_mean_squared_error: 57694.7930
Epoch 56/100
563/563  2s 3ms/step - loss: 3002795008.0000 - root_mean_squared_error: 54788.5547 - val_loss: 3369355776.0000 - val_root_mean_squared_error: 58046.1523
Epoch 57/100
563/563  2s 2ms/step - loss: 2948444416.0000 - root_mean_squared_error: 54295.2383 - val_loss: 3412904960.0000 - val_root_mean_squared_error: 58420.0742
Epoch 58/100
563/563  1s 2ms/step - loss: 2999850752.0000 - root_mean_squared_error: 54763.9961 - val_loss: 3373189632.0000 - val_root_mean_squared_error: 58079.1680
Epoch 59/100
563/563  1s 2ms/step - loss: 2920975616.0000 - root_mean_squared_error: 54038.0391 - val_loss: 3357112832.0000 - val_root_mean_squared_error: 57940.5977
Epoch 60/100
563/563  1s 2ms/step - loss: 2934187264.0000 - root_mean_squared_error: 54162.1406 - val_loss: 3397948672.0000 - val_root_mean_squared_error: 58291.9258
Epoch 61/100
563/563  1s 2ms/step - loss: 3040012288.0000 - root_mean_squared_error: 55123.9609 - val_loss: 3299003904.0000 - val_root_mean_squared_error: 57436.9570
Epoch 62/100
563/563  1s 2ms/step - loss: 2953530880.0000 - root_mean_squared_error: 54339.9414 - val_loss: 3324859648.0000 - val_root_mean_squared_error: 57661.5977
Epoch 63/100
563/563  1s 2ms/step - loss: 2907568896.0000 - root_mean_squared_error: 53917.2422 - val_loss: 3333906176.0000 - val_root_mean_squared_error: 57739.9883
Epoch 64/100
563/563  1s 2ms/step - loss: 2916231168.0000 - root_mean_squared_error: 53995.5977 - val_loss: 3450786048.0000 - val_root_mean_squared_error: 58743.3906
Epoch 65/100
563/563  2s 2ms/step - loss: 2947845120.0000 - root_mean_squared_error: 54275.5703 - val_loss: 3378517248.0000 - val_root_mean_squared_error: 58125.0156
Epoch 66/100
563/563  2s 2ms/step - loss: 2827737344.0000 - root_mean_squared_error: 53165.6797 - val_loss: 3598701312.0000 - val_root_mean_squared_error: 59989.1758
Epoch 67/100
563/563  1s 2ms/step - loss: 2910422784.0000 - root_mean_squared_error: 53943.3750 - val_loss: 3412234240.0000 - val_root_mean_squared_error: 58414.3320
Epoch 68/100
563/563  1s 2ms/step - loss: 2998081792.0000 - root_mean_squared_error: 54745.0117 - val_loss: 3675497728.0000 - val_root_mean_squared_error: 60625.8828
Epoch 69/100
563/563  2s 2ms/step - loss: 2865273856.0000 - root_mean_squared_error: 53506.3477 - val_loss: 3298195456.0000 - val_root_mean_squared_error: 57429.9180
Epoch 70/100
563/563  1s 2ms/step - loss: 3017387776.0000 - root_mean_squared_error: 54894.9062 - val_loss: 3316144384.0000 - val_root_mean_squared_error: 57585.9727
Epoch 71/100
563/563  1s 2ms/step - loss: 2856635648.0000 - root_mean_squared_error: 53506.3477 - val_loss: 3298195456.0000 - val_root_mean_squared_error: 57429.9180

r: 53440.0859 - val_loss: 3250820864.0000 - val_root_mean_squared_error: 57015.9688
Epoch 72/100
563/563  1s 2ms/step - loss: 2940389888.0000 - root_mean_squared_error: 54184.2344 - val_loss: 3554553856.0000 - val_root_mean_squared_error: 59620.0781
Epoch 73/100
563/563  1s 2ms/step - loss: 2920979456.0000 - root_mean_squared_error: 54036.2305 - val_loss: 3320546304.0000 - val_root_mean_squared_error: 57624.1797
Epoch 74/100
563/563  1s 2ms/step - loss: 2898277376.0000 - root_mean_squared_error: 53821.5703 - val_loss: 3397987840.0000 - val_root_mean_squared_error: 58292.2617
Epoch 75/100
563/563  2s 3ms/step - loss: 2936926464.0000 - root_mean_squared_error: 54170.9727 - val_loss: 3197381632.0000 - val_root_mean_squared_error: 56545.3945
Epoch 76/100
563/563  2s 3ms/step - loss: 2835487232.0000 - root_mean_squared_error: 53238.1719 - val_loss: 3338370048.0000 - val_root_mean_squared_error: 57778.6289
Epoch 77/100
563/563  2s 2ms/step - loss: 2803768064.0000 - root_mean_squared_error: 52944.4375 - val_loss: 3267643136.0000 - val_root_mean_squared_error: 57163.3008
Epoch 78/100
563/563  1s 2ms/step - loss: 2767419392.0000 - root_mean_squared_error: 52602.3125 - val_loss: 3323363584.0000 - val_root_mean_squared_error: 57648.6211
Epoch 79/100
563/563  1s 2ms/step - loss: 2731195904.0000 - root_mean_squared_error: 52252.7969 - val_loss: 3232067584.0000 - val_root_mean_squared_error: 56851.2773
Epoch 80/100
563/563  1s 2ms/step - loss: 2695593984.0000 - root_mean_squared_error: 51901.3164 - val_loss: 3266189824.0000 - val_root_mean_squared_error: 57150.5898
Epoch 81/100
563/563  1s 2ms/step - loss: 2820049152.0000 - root_mean_squared_error: 53094.2461 - val_loss: 3329429504.0000 - val_root_mean_squared_error: 57701.2070
Epoch 82/100
563/563  1s 2ms/step - loss: 2812450560.0000 - root_mean_squared_error: 53017.9141 - val_loss: 3207268608.0000 - val_root_mean_squared_error: 56632.7539
Epoch 83/100
563/563  1s 2ms/step - loss: 2801030912.0000 - root_mean_squared_error: 52918.9570 - val_loss: 3256188672.0000 - val_root_mean_squared_error: 57063.0234
Epoch 84/100
563/563  1s 2ms/step - loss: 2736392448.0000 - root_mean_squared_error: 52306.1055 - val_loss: 3294946816.0000 - val_root_mean_squared_error: 57401.6289
Epoch 85/100
563/563  2s 3ms/step - loss: 2802924800.0000 - root_mean_squared_error: 52934.9922 - val_loss: 3177144832.0000 - val_root_mean_squared_error: 56366.1680
Epoch 86/100
563/563  2s 3ms/step - loss: 2694107136.0000 - root_mean_squared_error: 51880.6758 - val_loss: 3183791616.0000 - val_root_mean_squared_error: 56425.0977
Epoch 87/100
563/563  2s 2ms/step - loss: 2705714176.0000 - root_mean_squared_error: 52010.1016 - val_loss: 3172370432.0000 - val_root_mean_squared_error: 56323.8008
Epoch 88/100
563/563  1s 2ms/step - loss: 2757981952.0000 - root_mean_squared_error: 52504.8828 - val_loss: 3556987904.0000 - val_root_mean_squared_error: 59640.4883
Epoch 89/100
563/563  1s 2ms/step - loss: 2719256576.0000 - root_mean_squared_error: 52134.2852 - val_loss: 3156543488.0000 - val_root_mean_squared_error: 56183.1250
Epoch 90/100
563/563  1s 2ms/step - loss: 2702067968.0000 - root_mean_squared_error: 51965.3047 - val_loss: 3149921536.0000 - val_root_mean_squared_error: 56124.1602
Epoch 91/100
563/563  1s 2ms/step - loss: 2615707648.0000 - root_mean_squared_error: 51122.2734 - val_loss: 3196517120.0000 - val_root_mean_squared_error: 56537.7500
Epoch 92/100
563/563  1s 2ms/step - loss: 2786473216.0000 - root_mean_squared_error: 52777.5391 - val_loss: 3376427776.0000 - val_root_mean_squared_error: 58107.0391
Epoch 93/100
563/563  1s 2ms/step - loss: 2750332672.0000 - root_mean_squared_error: 52439.9453 - val_loss: 3351586560.0000 - val_root_mean_squared_error: 57892.8906
Epoch 94/100
563/563  1s 2ms/step - loss: 2717035776.0000 - root_mean_squared_error: 52118.0703 - val_loss: 3342032640.0000 - val_root_mean_squared_error: 57810.3164
Epoch 95/100
563/563  2s 3ms/step - loss: 2710880768.0000 - root_mean_squared_error: 52010.1016 - val_loss: 3172370432.0000 - val_root_mean_squared_error: 56323.8008

```

r: 52056.3086 - val_loss: 3274188800.0000 - val_root_mean_squared_error: 57220.5273
Epoch 96/100
563/563 ————— 2s 2ms/step - loss: 2720423680.0000 - root_mean_squared_error: 52140.6211 - val_loss: 3125320960.0000 - val_root_mean_squared_error: 55904.5703
Epoch 97/100
563/563 ————— 1s 2ms/step - loss: 2712364288.0000 - root_mean_squared_error: 52068.3398 - val_loss: 3376818176.0000 - val_root_mean_squared_error: 58110.3945
Epoch 98/100
563/563 ————— 1s 2ms/step - loss: 2600215040.0000 - root_mean_squared_error: 50972.1562 - val_loss: 3128856320.0000 - val_root_mean_squared_error: 55936.1797
Epoch 99/100
563/563 ————— 1s 2ms/step - loss: 2713232128.0000 - root_mean_squared_error: 52086.4727 - val_loss: 3166954240.0000 - val_root_mean_squared_error: 56275.6992
Epoch 100/100
563/563 ————— 1s 2ms/step - loss: 2770647808.0000 - root_mean_squared_error: 52631.8164 - val_loss: 3049800192.0000 - val_root_mean_squared_error: 55224.9961

```

Out[]:

```
<keras.src.callbacks.history.History at 0x7880608257b0>
```

In []:

```

medium_nn = load_model('models/medium_nn.keras')
mse(medium_nn.predict(X_train), y_train, squared=False), mse(medium_nn.predict(X_val), y_val, squared=False)

```

```

563/563 ————— 1s 1ms/step
39/39 ————— 0s 1ms/step

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(

```

Out[]:

```
(52156.613, 55225.0)
```

3) for Large Neural Network

A highly complex network with multiple layers, useful for capturing intricate patterns.

In []:

```

large_nn = Sequential()
large_nn.add(InputLayer((13,)))
large_nn.add(Dense(256, 'relu'))
large_nn.add(Dense(128, 'relu'))
large_nn.add(Dense(64, 'relu'))
large_nn.add(Dense(32, 'relu'))
large_nn.add(Dense(1, 'linear'))

#Training Neural Network
#Neural networks require iterative optimization, where the optimizer (e.g., Adam) adjusts weights to minimize error.
opt = Adam(learning_rate=.1)
cp = ModelCheckpoint('models/large_nn.keras', save_best_only=True)
large_nn.compile(optimizer=opt, loss='mse', metrics=[RootMeanSquaredError()])
large_nn.fit(x=X_train, y=y_train, validation_data=(X_val, y_val), callbacks=[cp], epochs=100)

```

```

Epoch 1/100
563/563 ————— 3s 3ms/step - loss: 14577685504.0000 - root_mean_squared_error: 115639.1953 - val_loss: 6201869312.0000 - val_root_mean_squared_error: 78751.9453
Epoch 2/100
563/563 ————— 2s 2ms/step - loss: 4583801344.0000 - root_mean_squared_error: 67686.6641 - val_loss: 4813209600.0000 - val_root_mean_squared_error: 69377.2969
Epoch 3/100

```


563/563 ————— 2s 4ms/step - loss: 4220409344.0000 - root_mean_squared_error: 64951.8555 - val_loss: 5594952704.0000 - val_root_mean_squared_error: 74799.4141
Epoch 4/100

563/563 ————— 2s 3ms/step - loss: 4146496256.0000 - root_mean_squared_error: 64378.2812 - val_loss: 4005124608.0000 - val_root_mean_squared_error: 63286.0547
Epoch 5/100

563/563 ————— 1s 2ms/step - loss: 3633795328.0000 - root_mean_squared_error: 60272.2188 - val_loss: 4249569536.0000 - val_root_mean_squared_error: 65188.7227
Epoch 6/100

563/563 ————— 1s 2ms/step - loss: 3642889728.0000 - root_mean_squared_error: 60349.5273 - val_loss: 3903959808.0000 - val_root_mean_squared_error: 62481.6758
Epoch 7/100

563/563 ————— 1s 2ms/step - loss: 3608773120.0000 - root_mean_squared_error: 60055.6719 - val_loss: 3891719680.0000 - val_root_mean_squared_error: 62383.6484
Epoch 8/100

563/563 ————— 3s 2ms/step - loss: 3386950144.0000 - root_mean_squared_error: 58176.4805 - val_loss: 3559396864.0000 - val_root_mean_squared_error: 59660.6797
Epoch 9/100

563/563 ————— 1s 2ms/step - loss: 3472865280.0000 - root_mean_squared_error: 58906.1953 - val_loss: 3635167232.0000 - val_root_mean_squared_error: 60292.3477
Epoch 10/100

563/563 ————— 1s 3ms/step - loss: 3272971520.0000 - root_mean_squared_error: 57186.0938 - val_loss: 4142724096.0000 - val_root_mean_squared_error: 64363.9961
Epoch 11/100

563/563 ————— 2s 4ms/step - loss: 3476674304.0000 - root_mean_squared_error: 58941.3984 - val_loss: 4091129856.0000 - val_root_mean_squared_error: 63961.9414
Epoch 12/100

563/563 ————— 2s 3ms/step - loss: 3373266432.0000 - root_mean_squared_error: 58050.8008 - val_loss: 3671097344.0000 - val_root_mean_squared_error: 60589.5820
Epoch 13/100

563/563 ————— 1s 2ms/step - loss: 3230081792.0000 - root_mean_squared_error: 56827.1680 - val_loss: 3720891648.0000 - val_root_mean_squared_error: 60999.1133
Epoch 14/100

563/563 ————— 2s 2ms/step - loss: 3200335104.0000 - root_mean_squared_error: 56560.8555 - val_loss: 3624494592.0000 - val_root_mean_squared_error: 60203.7773
Epoch 15/100

563/563 ————— 1s 2ms/step - loss: 3000845568.0000 - root_mean_squared_error: 54776.1641 - val_loss: 3475805440.0000 - val_root_mean_squared_error: 58955.9609
Epoch 16/100

563/563 ————— 2s 2ms/step - loss: 3199978496.0000 - root_mean_squared_error: 56555.0469 - val_loss: 3342490624.0000 - val_root_mean_squared_error: 57814.2773
Epoch 17/100

563/563 ————— 3s 3ms/step - loss: 3133114624.0000 - root_mean_squared_error: 55958.9961 - val_loss: 3940275200.0000 - val_root_mean_squared_error: 62771.6133
Epoch 18/100

563/563 ————— 2s 3ms/step - loss: 2926949632.0000 - root_mean_squared_error: 54083.2812 - val_loss: 3205490176.0000 - val_root_mean_squared_error: 56617.0469
Epoch 19/100

563/563 ————— 2s 2ms/step - loss: 2936738560.0000 - root_mean_squared_error: 54177.9844 - val_loss: 3238283776.0000 - val_root_mean_squared_error: 56905.9219
Epoch 20/100

563/563 ————— 1s 2ms/step - loss: 2983696896.0000 - root_mean_squared_error: 54601.5195 - val_loss: 3240989696.0000 - val_root_mean_squared_error: 56929.6914
Epoch 21/100

563/563 ————— 1s 2ms/step - loss: 2800329216.0000 - root_mean_squared_error: 52897.8438 - val_loss: 3930374144.0000 - val_root_mean_squared_error: 62692.6953
Epoch 22/100

563/563 ————— 3s 2ms/step - loss: 2830406656.0000 - root_mean_squared_error: 53186.4414 - val_loss: 3484327936.0000 - val_root_mean_squared_error: 59028.1953
Epoch 23/100

563/563 ————— 1s 2ms/step - loss: 2929417472.0000 - root_mean_squared_error: 54115.7500 - val_loss: 3241837568.0000 - val_root_mean_squared_error: 56937.1367
Epoch 24/100

563/563 ————— 2s 3ms/step - loss: 2801490432.0000 - root_mean_squared_error: 52905.1016 - val_loss: 3053377536.0000 - val_root_mean_squared_error: 55257.3750
Epoch 25/100

563/563 ————— 2s 4ms/step - loss: 2815137280.0000 - root_mean_squared_error: 53031.0039 - val_loss: 3200846080.0000 - val_root_mean_squared_error: 56576.0195
Epoch 26/100

563/563 ————— 2s 3ms/step - loss: 2819328000.0000 - root_mean_squared_error: 53092.1250 - val_loss: 3417561088.0000 - val_root_mean_squared_error: 58459.9102
Epoch 27/100

563/563 ————— 2s 2ms/step - loss: 2726735360.0000 - root_mean_squared_error: 52211.3008 - val_loss: 3274870016.0000 - val_root_mean_squared_error: 57226.4805
Epoch 28/100

563/563 ————— 1s 2ms/step - loss: 2852363264.0000 - root_mean_squared_error: 53390.0195 - val_loss: 3414737920.0000 - val_root_mean_squared_error: 58435.7578
Epoch 29/100

563/563 ————— 1s 2ms/step - loss: 2770728192.0000 - root_mean_squared_error: 52624.4492 - val_loss: 3434776576.0000 - val_root_mean_squared_error: 58606.9688
Epoch 30/100

563/563 ————— 1s 2ms/step - loss: 2789532160.0000 - root_mean_squared_error: 52802.5547 - val_loss: 3262315520.0000 - val_root_mean_squared_error: 57116.6836
Epoch 31/100

563/563 ————— 3s 2ms/step - loss: 2683632896.0000 - root_mean_squared_error: 51798.2422 - val_loss: 3737213440.0000 - val_root_mean_squared_error: 61132.7539
Epoch 32/100

563/563 ————— 2s 4ms/step - loss: 2675304192.0000 - root_mean_squared_error: 51712.0391 - val_loss: 3171623168.0000 - val_root_mean_squared_error: 56317.1641
Epoch 33/100

563/563 ————— 2s 4ms/step - loss: 2906547200.0000 - root_mean_squared_error: 53903.5508 - val_loss: 3870891520.0000 - val_root_mean_squared_error: 62216.4883
Epoch 34/100

563/563 ————— 2s 2ms/step - loss: 2692177920.0000 - root_mean_squared_error: 51868.2695 - val_loss: 3158903296.0000 - val_root_mean_squared_error: 56204.1211
Epoch 35/100

563/563 ————— 1s 2ms/step - loss: 2738666496.0000 - root_mean_squared_error: 52324.6719 - val_loss: 3229836288.0000 - val_root_mean_squared_error: 56831.6484
Epoch 36/100

563/563 ————— 2s 2ms/step - loss: 2649613568.0000 - root_mean_squared_error: 51456.8047 - val_loss: 3121056768.0000 - val_root_mean_squared_error: 55866.4180
Epoch 37/100

563/563 ————— 1s 2ms/step - loss: 2725374464.0000 - root_mean_squared_error: 52188.0586 - val_loss: 3146051840.0000 - val_root_mean_squared_error: 56089.6758
Epoch 38/100

563/563 ————— 1s 2ms/step - loss: 2729904640.0000 - root_mean_squared_error: 52244.3594 - val_loss: 3123862528.0000 - val_root_mean_squared_error: 55891.5234
Epoch 39/100

563/563 ————— 3s 3ms/step - loss: 2681320192.0000 - root_mean_squared_error: 51773.8789 - val_loss: 3806076160.0000 - val_root_mean_squared_error: 61693.4062
Epoch 40/100

563/563 ————— 2s 4ms/step - loss: 2695783424.0000 - root_mean_squared_error: 51904.8789 - val_loss: 3033415680.0000 - val_root_mean_squared_error: 55076.4531
Epoch 41/100

563/563 ————— 1s 2ms/step - loss: 2676339200.0000 - root_mean_squared_error: 51722.9922 - val_loss: 4016765696.0000 - val_root_mean_squared_error: 63377.9609
Epoch 42/100

563/563 ————— 3s 2ms/step - loss: 2690556160.0000 - root_mean_squared_error: 51855.9062 - val_loss: 3145768448.0000 - val_root_mean_squared_error: 56087.1523
Epoch 43/100

563/563 ————— 3s 2ms/step - loss: 2669172480.0000 - root_mean_squared_error: 51655.9688 - val_loss: 3064726016.0000 - val_root_mean_squared_error: 55359.9688
Epoch 44/100

563/563 ————— 1s 2ms/step - loss: 2628993536.0000 - root_mean_squared_error: 51267.3164 - val_loss: 3216302592.0000 - val_root_mean_squared_error: 56712.4570
Epoch 45/100

563/563 ————— 1s 2ms/step - loss: 2689366016.0000 - root_mean_squared_error: 51845.2773 - val_loss: 3065277952.0000 - val_root_mean_squared_error: 55364.9531
Epoch 46/100

563/563 ————— 2s 3ms/step - loss: 2628811776.0000 - root_mean_squared_error: 51264.6992 - val_loss: 3089552896.0000 - val_root_mean_squared_error: 55583.7461
Epoch 47/100

563/563 ————— 2s 4ms/step - loss: 2566642688.0000 - root_mean_squared_error: 50640.9062 - val_loss: 3068019200.0000 - val_root_mean_squared_error: 55389.7031
Epoch 48/100

563/563 ————— 2s 4ms/step - loss: 2688141312.0000 - root_mean_squared_error: 51812.4023 - val_loss: 3369011968.0000 - val_root_mean_squared_error: 58043.1914
Epoch 49/100

563/563 ————— 2s 2ms/step - loss: 2676328960.0000 - root_mean_squared_error: 51730.8008 - val_loss: 2920869888.0000 - val_root_mean_squared_error: 54045.0742
Epoch 50/100

563/563 ————— 1s 2ms/step - loss: 2496260864.0000 - root_mean_squared_error: 49948.8711 - val_loss: 3018246912.0000 - val_root_mean_squared_error: 54938.5742
Epoch 51/100

563/563 ————— 3s 2ms/step - loss: 2580494336.0000 - root_mean_squared_error: 50791.3164 - val_loss: 2951595264.0000 - val_root_mean_squared_error: 54328.5859
Epoch 52/100

563/563 ————— 1s 2ms/step - loss: 2675419392.0000 - root_mean_squared_error: 51707.8516 - val_loss: 3352327424.0000 - val_root_mean_squared_error: 57899.2852
Epoch 53/100

563/563 ————— 1s 2ms/step - loss: 2552685824.0000 - root_mean_squared_error: 50503.9609 - val_loss: 3785790976.0000 - val_root_mean_squared_error: 61528.7812
Epoch 54/100

563/563 ————— 4s 4ms/step - loss: 2557549824.0000 - root_mean_squared_error: 50567.8008 - val_loss: 3115501568.0000 - val_root_mean_squared_error: 55816.6797
Epoch 55/100

563/563 ————— 3s 4ms/step - loss: 2514075392.0000 - root_mean_squared_error: 50124.4102 - val_loss: 2995494144.0000 - val_root_mean_squared_error: 54731.1094
Epoch 56/100

563/563 ————— 2s 4ms/step - loss: 2553985536.0000 - root_mean_squared_error: 50529.9766 - val_loss: 3158546688.0000 - val_root_mean_squared_error: 56200.9492
Epoch 57/100

563/563 ————— 1s 2ms/step - loss: 2590243584.0000 - root_mean_squared_error: 50872.2852 - val_loss: 3221155072.0000 - val_root_mean_squared_error: 56755.2188
Epoch 58/100

563/563 ————— 3s 2ms/step - loss: 2505511680.0000 - root_mean_squared_error: 50037.0547 - val_loss: 3064176384.0000 - val_root_mean_squared_error: 55355.0039
Epoch 59/100

563/563 ————— 1s 2ms/step - loss: 2668993792.0000 - root_mean_squared_error: 51649.0625 - val_loss: 3111876096.0000 - val_root_mean_squared_error: 55784.1914
Epoch 60/100

563/563 ————— 3s 2ms/step - loss: 2455049728.0000 - root_mean_squared_error: 49526.9258 - val_loss: 3074504704.0000 - val_root_mean_squared_error: 55448.2148
Epoch 61/100

563/563 ————— 2s 3ms/step - loss: 2616104192.0000 - root_mean_squared_error: 51141.8477 - val_loss: 3059302912.0000 - val_root_mean_squared_error: 55310.9648
Epoch 62/100

563/563 ————— 3s 4ms/step - loss: 2450646528.0000 - root_mean_squared_error: 49488.2734 - val_loss: 3288945152.0000 - val_root_mean_squared_error: 57349.3242
Epoch 63/100

563/563 ————— 2s 2ms/step - loss: 2582804992.0000 - root_mean_squared_error: 50816.5664 - val_loss: 3403234560.0000 - val_root_mean_squared_error: 58337.2500
Epoch 64/100

563/563 ————— 1s 2ms/step - loss: 2603771904.0000 - root_mean_squared_error: 51014.6523 - val_loss: 3203316736.0000 - val_root_mean_squared_error: 56597.8516
Epoch 65/100

563/563 ————— 3s 3ms/step - loss: 2639634944.0000 - root_mean_squared_error: 51366.7812 - val_loss: 3038631168.0000 - val_root_mean_squared_error: 55123.7812
Epoch 66/100

563/563 ————— 1s 2ms/step - loss: 2516060160.0000 - root_mean_squared_error: 50153.4609 - val_loss: 2941582336.0000 - val_root_mean_squared_error: 54236.3555
Epoch 67/100

563/563 ————— 1s 2ms/step - loss: 2492644864.0000 - root_mean_squared_error: 49915.3711 - val_loss: 3040054016.0000 - val_root_mean_squared_error: 55136.6836
Epoch 68/100

563/563 ————— 3s 4ms/step - loss: 2482366208.0000 - root_mean_squared_error: 49817.3281 - val_loss: 3214752512.0000 - val_root_mean_squared_error: 56698.7891
Epoch 69/100

563/563 ————— 2s 3ms/step - loss: 2425843968.0000 - root_mean_squared_error: 49240.5664 - val_loss: 3042396672.0000 - val_root_mean_squared_error: 55157.9258
Epoch 70/100

563/563 ————— 2s 2ms/step - loss: 2462455552.0000 - root_mean_squared_error: 49595.4336 - val_loss: 3221079552.0000 - val_root_mean_squared_error: 56754.5547
Epoch 71/100

563/563 ————— 2s 2ms/step - loss: 2454253312.0000 - root_mean_squared_error: 49530.9531 - val_loss: 3042770944.0000 - val_root_mean_squared_error: 55161.3164
Epoch 72/100

563/563 ————— 1s 2ms/step - loss: 2423707904.0000 - root_mean_squared_error: 49211.4297 - val_loss: 3126023168.0000 - val_root_mean_squared_error: 55910.8516
Epoch 73/100

563/563 ————— 3s 2ms/step - loss: 2559433984.0000 - root_mean_squared_error: 50567.7109 - val_loss: 3018620160.0000 - val_root_mean_squared_error: 54941.9727
Epoch 74/100

563/563 ————— 3s 4ms/step - loss: 2611820544.0000 - root_mean_squared_error: 51095.9570 - val_loss: 3227091200.0000 - val_root_mean_squared_error: 56807.4922
Epoch 75/100

563/563 ————— 1s 3ms/step - loss: 2554323456.0000 - root_mean_squared_error: 50532.6523 - val_loss: 2990388480.0000 - val_root_mean_squared_error: 54684.4453
Epoch 76/100

563/563 ————— 1s 2ms/step - loss: 2524417024.0000 - root_mean_squared_error: 50216.8281 - val_loss: 3016230144.0000 - val_root_mean_squared_error: 54920.2148
Epoch 77/100

563/563 ————— 1s 2ms/step - loss: 2562827264.0000 - root_mean_squared_error: 50601.2383 - val_loss: 3110782208.0000 - val_root_mean_squared_error: 55774.3867
Epoch 78/100

563/563 ————— 3s 2ms/step - loss: 2510818048.0000 - root_mean_squared_error: 50095.9805 - val_loss: 3107752704.0000 - val_root_mean_squared_error: 55747.2227
Epoch 79/100

563/563 ————— 1s 2ms/step - loss: 2418304512.0000 - root_mean_squared_error: 49171.0078 - val_loss: 2953392896.0000 - val_root_mean_squared_error: 54345.1289
Epoch 80/100

563/563 ————— 1s 2ms/step - loss: 2422451456.0000 - root_mean_squared_error: 49216.2734 - val_loss: 3055053312.0000 - val_root_mean_squared_error: 55272.5352
Epoch 81/100

563/563 ————— 3s 4ms/step - loss: 2455007488.0000 - root_mean_squared_error: 49531.5781 - val_loss: 3045969408.0000 - val_root_mean_squared_error: 55190.3008
Epoch 82/100

563/563 ————— 2s 2ms/step - loss: 2419768576.0000 - root_mean_squared_error: 49178.2539 - val_loss: 3074956544.0000 - val_root_mean_squared_error: 55452.2891
Epoch 83/100

563/563 ————— 1s 2ms/step - loss: 2475449856.0000 - root_mean_squared_error: 49741.2188 - val_loss: 3053578240.0000 - val_root_mean_squared_error: 55259.1914
Epoch 84/100

563/563 ————— 3s 2ms/step - loss: 2460014848.0000 - root_mean_squared_error: 49586.3242 - val_loss: 2930732032.0000 - val_root_mean_squared_error: 54136.2344
Epoch 85/100

563/563 ————— 1s 2ms/step - loss: 2466914816.0000 - root_mean_squared_error: 49663.2266 - val_loss: 3089690624.0000 - val_root_mean_squared_error: 55584.9844
Epoch 86/100

563/563 ————— 1s 2ms/step - loss: 2354352896.0000 - root_mean_squared_error: 48514.3359 - val_loss: 2972948480.0000 - val_root_mean_squared_error: 54524.7500
Epoch 87/100

563/563 ————— 1s 2ms/step - loss: 2393474560.0000 - root_mean_squared_error: 48903.1094 - val_loss: 3066971392.0000 - val_root_mean_squared_error: 55380.2422
Epoch 88/100

563/563 ————— 3s 4ms/step - loss: 2507817728.0000 - root_mean_squared_error: 50065.5938 - val_loss: 3044221184.0000 - val_root_mean_squared_error: 55174.4609
Epoch 89/100

563/563 ————— 2s 2ms/step - loss: 2438450944.0000 - root_mean_squared_error: 49375.7148 - val_loss: 2939404288.0000 - val_root_mean_squared_error: 54216.2734
Epoch 90/100

563/563 ————— 3s 2ms/step - loss: 2373002752.0000 - root_mean_squared_error: 48706.0273 - val_loss: 3084823296.0000 - val_root_mean_squared_error: 55541.1875
Epoch 91/100

563/563 ————— 3s 2ms/step - loss: 2367094272.0000 - root_mean_squared_error: 48645.6484 - val_loss: 3107561216.0000 - val_root_mean_squared_error: 55745.5039
Epoch 92/100

563/563 ————— 1s 2ms/step - loss: 2442265344.0000 - root_mean_squared_error: 49413.9883 - val_loss: 3187333632.0000 - val_root_mean_squared_error: 56456.4766
Epoch 93/100

563/563 ————— 1s 2ms/step - loss: 2420281088.0000 - root_mean_squared_error: 49183.8438 - val_loss: 3003081728.0000 - val_root_mean_squared_error: 54800.3789
Epoch 94/100

563/563 ————— 2s 3ms/step - loss: 2492461312.0000 - root_mean_squared_error: 49917.5781 - val_loss: 2895421952.0000 - val_root_mean_squared_error: 53809.1250
Epoch 95/100

563/563 ————— 2s 4ms/step - loss: 2491831296.0000 - root_mean_squared_error: 49912.1211 - val_loss: 3070292992.0000 - val_root_mean_squared_error: 55410.2266
Epoch 96/100

563/563 ————— 2s 2ms/step - loss: 2385057024.0000 - root_mean_squared_error: 48829.8594 - val_loss: 3248151296.0000 - val_root_mean_squared_error: 56992.5547
Epoch 97/100

563/563 ————— 2s 2ms/step - loss: 2567180032.0000 - root_mean_squared_error: 50659.0664 - val_loss: 3005500416.0000 - val_root_mean_squared_error: 54822.4453
Epoch 98/100

563/563 ————— 2s 2ms/step - loss: 2410853632.0000 - root_mean_squared_error: 49095.7734 - val_loss: 3089205248.0000 - val_root_mean_squared_error: 55580.6211
Epoch 99/100

```
563/563 ————— 1s 2ms/step - loss: 2425818880.0000 - root_mean_squared_error: 49246.5039 - val_loss: 3288166400.0000 - val_root_mean_squared_error: 57342.5352
Epoch 100/100
563/563 ————— 1s 2ms/step - loss: 2372902144.0000 - root_mean_squared_error: 48691.7617 - val_loss: 2922787584.0000 - val_root_mean_squared_error: 54062.8125
```

Out[]:

```
<keras.src.callbacks.history.History at 0x788062620e80>
```

In []:

```
large_nn = load_model('models/large_nn.keras')
mse(large_nn.predict(X_train), y_train, squared=False), mse(large_nn.predict(X_val), y_val, squared=False)
```

```
563/563 ————— 1s 1ms/step
39/39 ————— 0s 1ms/step
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

Out[]:

```
(47861.855, 53809.13)
```

In []:

```
mse(gbr.predict(X_test), y_test, squared=False)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning
: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
  warnings.warn(
```

Out[]:

```
49345.935411338134
```

Evaluation

Models are evaluated using the Root Mean Squared Error (RMSE):

1) Linear Regression: Poor performance due to limited complexity.

2) KNN: Improved accuracy but risks overfitting with fewer neighbors.

3) Random Forest: Strong performance with manageable complexity.

4) Gradient Boosting: Best among traditional models, achieving low RMSE.

5) Neural Networks: Increasing complexity improves accuracy, but the risk of overfitting rises with deeper architectures.

CONCLUSION

The most suitable Model here is Gradient Boosting achieved the best overall performance on the test set with the lowest RMSE (~49,000).

while Simpler models (like Linear Regression) are less effective for complex datasets.

Ensemble methods like Random Forest and Gradient Boosting excel in capturing non-linear relationships. and

Neural networks require careful tuning to balance complexity and overfitting. This project highlights the importance of trying multiple models, preprocessing data effectively, and evaluating results to make informed decisions.

