

The project aims to predict the selling price of used cars using machine learning techniques. We'll preprocess the dataset, encode categorical variables, train two models (Linear Regression and Lasso Regression), and evaluate their performance. Additionally, we will use visualizations to compare actual and predicted prices to understand model performance better.

In [1]:

```
#IMPORTING NECESSARY DEPENDENCIES AND LIBRARY
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso
from sklearn import metrics
```

In [2]:

```
#LOADING THE DATASET
dataset = pd.read_csv('/content/car_data.csv')
```

In [5]:

```
#CHECKING FIRST AND LAST FIVE ROWS OF DATASET

dataset.head()
```

Out[5]:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
296	city	2016	9.50	11.6	33988	Diesel	Dealer	Manual	0
297	brio	2015	4.00	5.9	60000	Petrol	Dealer	Manual	0
298	city	2009	3.35	11.0	87934	Petrol	Dealer	Manual	0
299	city	2017	11.50	12.5	9000	Diesel	Dealer	Manual	0
300	brio	2016	5.30	5.9	5464	Petrol	Dealer	Manual	0

In [6]:

```
dataset.tail()
```

Out[6]:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
296	city	2016	9.50	11.6	33988	Diesel	Dealer	Manual	0
297	brio	2015	4.00	5.9	60000	Petrol	Dealer	Manual	0
298	city	2009	3.35	11.0	87934	Petrol	Dealer	Manual	0
299	city	2017	11.50	12.5	9000	Diesel	Dealer	Manual	0
300	brio	2016	5.30	5.9	5464	Petrol	Dealer	Manual	0

In [8]:

```
#NUMBER OF ROWS AND COLUMN
dataset.shape
```

Out[8]:

```
(301, 9)
```

In [9]:

# #INFORMATION ABOUT DATASET

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Car_Name              301 non-null   object
1   Year                  301 non-null   int64
2   Selling_Price         301 non-null   float64
3   Present_Price         301 non-null   float64
4   Kms_Driven            301 non-null   int64
5   Fuel_Type             301 non-null   object
6   Seller_Type           301 non-null   object
7   Transmission          301 non-null   object
8   Owner                 301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

In [11]:

## #CHECKING NULL VALUES

```
print("\nMissing values in each column:")
dataset.isnull().sum()
```

Missing values in each column:

Out[11]:

	0
Car_Name	0
Year	0
Selling_Price	0
Present_Price	0
Kms_Driven	0
Fuel_Type	0
Seller_Type	0
Transmission	0
Owner	0

dtype: int64

In [12]:

## #CHECKING THE DISTRIBUTION OF CATEGORICAL DATA

```
print("\nDistribution of Categorical Data:")
print("Seller_Type:", dataset.Seller_Type.value_counts())
print("Transmission:", dataset.Transmission.value_counts())
print("Fuel_Type:", dataset.Fuel_Type.value_counts())
```

Distribution of Categorical Data:

Seller\_Type: Seller\_Type

Dealer 195

Individual 106

Name: count, dtype: int64

Transmission: Transmission

Manual 261

Automatic 40

Name: count, dtype: int64

Fuel\_Type: Fuel\_Type

Petrol 239

Diesel 60

CNG 2

Name: count, dtype: int64

In [13]:

```
#ENCODING THE CATEGORICAL VARIABLES
```

```
dataset.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2}}, inplace=True)
dataset.replace({'Seller_Type': {'Dealer': 0, 'Individual': 1}}, inplace=True)
dataset.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)
```

```
<ipython-input-13-cdad747965b7>:2: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
```

```
dataset.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2}}, inplace=True)
```

```
<ipython-input-13-cdad747965b7>:3: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
```

```
dataset.replace({'Seller_Type': {'Dealer': 0, 'Individual': 1}}, inplace=True)
```

```
<ipython-input-13-cdad747965b7>:4: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
```

```
dataset.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)
```

In [14]:

```
#DISPLAYING THE UPDATED DATASET
```

```
print(dataset.head())
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	\
0	ritz	2014	3.35	5.59	27000	0	
1	sx4	2013	4.75	9.54	43000	1	
2	ciaz	2017	7.25	9.85	6900	0	
3	wagon r	2011	2.85	4.15	5200	0	
4	swift	2014	4.60	6.87	42450	1	

	Seller_Type	Transmission	Owner
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

In [15]:

```
# Splitting the data into features (X) and target (Y)
```

```
X = dataset.drop(['Car_Name', 'Selling_Price'], axis=1)
```

```
Y = dataset['Selling_Price']
```

In [16]:

```
print("\nFeature matrix (X):")
print(X.head())
print("\nTarget variable (Y):")
print(Y.head())
```

Feature matrix (X):

	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	\
0	2014	5.59	27000	0	0	0	
1	2013	9.54	43000	1	0	0	
2	2017	9.85	6900	0	0	0	
3	2011	4.15	5200	0	0	0	
4	2014	6.87	42450	1	0	0	

	Owner
0	0
1	0
2	0
3	0
4	0

Target variable (Y):

```
0    3.35
1    4.75
2    7.25
3    2.85
4    4.60
Name: Selling_Price, dtype: float64
```

In [17]:

```
#SPLITTING THE DATA INTO TRAINING AND TESTING SET
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

In [18]:

```
# Displaying dataset shapes
print("\nShapes of datasets:")
print("X_train:", X_train.shape, "X_test:", X_test.shape)
```

Shapes of datasets:  
X\_train: (240, 7) X\_test: (61, 7)

In [20]:

```
#Function to visualize actual vs predicted prices
def visualize_predictions(actual, predicted, title):
    plt.figure(figsize=(8, 6))
    sns.scatterplot(x=actual, y=predicted, alpha=0.6, edgecolor='b')
    plt.plot([min(actual), max(actual)], [min(actual), max(actual)], color='red', linewidth=2)
    plt.xlabel("Actual Price")
    plt.ylabel("Predicted Price")
    plt.title(title)
    plt.show()
```

## MODEL1 : LINEAR REGRESSION

In [21]:

```
print("\n--- Linear Regression Model ---")
lin_reg_model = LinearRegression()
lin_reg_model.fit(X_train, Y_train)
```

--- Linear Regression Model ---

Out[21]:

▼ LinearRegression ⓘ ?

LinearRegression()

In [22]:

```
# Evaluating Linear Regression model
train_pred_lr = lin_reg_model.predict(X_train)
test_pred_lr = lin_reg_model.predict(X_test)
```

In [23]:

```
train_r2_lr = metrics.r2_score(Y_train, train_pred_lr)
test_r2_lr = metrics.r2_score(Y_test, test_pred_lr)
```

In [24]:

```
print("Training R-squared score:", train_r2_lr)
print("Testing R-squared score:", test_r2_lr)
```


Training R-squared score: 0.8838169193709796  
Testing R-squared score: 0.8401532365377697

In [25]:


21 [10]:

```
# Visualizing actual vs predicted prices for Linear Regression
visualize_predictions(Y_train, train_pred_lr, "Linear Regression - Training Data")
visualize_predictions(Y_test, test_pred_lr, "Linear Regression - Testing Data")
```

Linear Regression - Training Data



Linear Regression - Testing Data



## MODEL 2 : LASSO REGRESSION

In [26]:

```
print("\n--- Lasso Regression Model ---")
lass_reg_model = Lasso()
lass_reg_model.fit(X_train, Y_train)
```

--- Lasso Regression Model ---

Out[26]:

▼ Lasso i ?

Lasso()

In [27]:

```
# Evaluating Lasso Regression model
train_pred_lasso = lass_reg_model.predict(X_train)
test_pred_lasso = lass_reg_model.predict(X_test)
```

In [28]:

```
train_r2_lasso = metrics.r2_score(Y_train, train_pred_lasso)
test_r2_lasso = metrics.r2_score(Y_test, test_pred_lasso)
```

In [29]:

```
print("Training R-squared score:", train_r2_lasso)
print("Testing R-squared score:", test_r2_lasso)
```


Training R-squared score: 0.843587395258283

Testing R-squared score: 0.8497457570738539

In [30]:

```
# Visualizing actual vs predicted prices for Lasso Regression
visualize_predictions(Y_train, train_pred_lasso, "Lasso Regression - Training Data")
visualize_predictions(Y_test, test_pred_lasso, "Lasso Regression - Testing Data")
```

Lasso Regression - Training Data



0 5 10 15 20 25 30 35  
Actual Price

Lasso Regression - Testing Data

